

# APSC-5984 Lab 6: Tidy Data

---

Due: 2023-02-27 (Monday) 23:59:59

- [APSC-5984 Lab 6: Tidy Data](#)
  - [0. Overview](#)
  - [1. Column headers are values, not variable names](#)
    - [1.1 United States Census Bureau: 2017 National Population Projection](#)
    - [1.2 Billboard top hits for 2000](#)
  - [2. Multiple variables are stored in one column](#)
    - [2.1 Tuberculosis \(TB\) dataset](#)
  - [3. Variables are stored in both rows and columns](#)
  - [4. Multiple types of observational units are stored in the same table](#)
    - [Get the columns that contain only song information](#)
    - [Get the columns that contain only song information](#)
  - [5. A single observational unit is stored in multiple tables](#)
  - [6. Case study](#)
    - [6.1 Basic data exploration](#)
    - [6.2 Count the number of deaths in each hour](#)
    - [6.3 The proportion of deaths in each cause by hour](#)
    - [6.4 The proportion of deaths in each hour](#)
    - [6.5 Deviation from the expected proportion](#)
    - [6.6 Use a residual plot to find unusual temporal patterns](#)
    - [6.7 Visualize the unusual temporal patterns](#)

## 0. Overview

Tidy data is an important concept in data science. We will go through each type of messy data discussed in the paper [Tidy Data](#) and learn how to tidy them. We will not only use the [pandas](#) library to manipulate dataframes, but also use [plotnine](#), which is a Python implementation of the [ggplot2](#) library in R, to visualize the tidy data.

First, let's import the libraries we will use in this section.

```
import pandas as pd
import numpy as np
import re
import os
from plotnine import *
```

## 1. Column headers are values, not variable names

### 1.1 United States Census Bureau: 2017 National Population Projection

This is a dataset describing the projected population of the United States from 2016 to 2060. The first column **YEAR** keeps the year of the project. The remaining columns **POP\_X** were used to store the projected population of each age **X** in the year, where **X** is an integer from 0 to 100.

```
data = pd.read_csv("tidy_1_pop.csv")
data
```

|   | YEAR | POP_0   | POP_1   | POP_2   | POP_3   | POP_4   | POP_5   | POP_6   | \ |
|---|------|---------|---------|---------|---------|---------|---------|---------|---|
| 0 | 2016 | 3970145 | 3995008 | 3992154 | 3982074 | 3987656 | 4032515 | 4029655 |   |
| 1 | 2017 | 4054035 | 3982964 | 4008116 | 4003478 | 3992207 | 3997392 | 4042440 |   |
| 2 | 2018 | 4075563 | 4068172 | 3995888 | 4019345 | 4013649 | 4001995 | 4007421 |   |
| 3 | 2019 | 4095614 | 4089881 | 4082231 | 4006967 | 4029427 | 4023461 | 4012057 |   |
| 4 | 2020 | 4113164 | 4110117 | 4104058 | 4094281 | 4016919 | 4039164 | 4033531 |   |

|          | POP_7   | POP_8   | ... | POP_91 | POP_92 | POP_93 | POP_94 | POP_95 |  |
|----------|---------|---------|-----|--------|--------|--------|--------|--------|--|
| POP_96 \ |         |         |     |        |        |        |        |        |  |
| 0        | 4029991 | 4159114 | ... | 449986 | 372625 | 300000 | 239313 | 186408 |  |
| 1        | 4040047 | 4041063 | ... | 449945 | 382669 | 311525 | 246219 | 192531 |  |
| 2        | 4052927 | 4051175 | ... | 462335 | 382993 | 320285 | 256011 | 198354 |  |
| 3        | 4017972 | 4064123 | ... | 467488 | 393919 | 320884 | 263533 | 206526 |  |
| 4        | 4022626 | 4029209 | ... | 464985 | 398712 | 330389 | 264318 | 212880 |  |

|   | POP_97 | POP_98 | POP_99 | POP_100 |
|---|--------|--------|--------|---------|
| 0 | 94311  | 68972  | 44895  | 81896   |
| 1 | 104540 | 70840  | 50486  | 83574   |
| 2 | 113165 | 78659  | 51938  | 86221   |
| 3 | 117240 | 85265  | 57778  | 87671   |
| 4 | 121128 | 88491  | 62724  | 92064   |

[45 rows x 102 columns]

We can use `pd.melt()` to tidy the dataset. The `id_vars` argument specifies the columns that should not be melted. The `value_vars` argument specifies the columns that should be melted. In our case, **YEAR** is `id_vars` and the remaining columns are `value_vars`. We can specify the names of the melted columns using the `var_name` and `value_name` arguments.

```
data_long = pd.melt(data,
                    id_vars=["YEAR"],
                    var_name="age",
                    value_name="pop")
data_long
```

|      | YEAR | age     | pop     |
|------|------|---------|---------|
| 0    | 2016 | POP_0   | 3970145 |
| 1    | 2017 | POP_0   | 4054035 |
| 2    | 2018 | POP_0   | 4075563 |
| 3    | 2019 | POP_0   | 4095614 |
| 4    | 2020 | POP_0   | 4113164 |
| ...  | ...  | ...     | ...     |
| 4540 | 2056 | POP_100 | 505951  |
| 4541 | 2057 | POP_100 | 529280  |
| 4542 | 2058 | POP_100 | 549748  |
| 4543 | 2059 | POP_100 | 567379  |
| 4544 | 2060 | POP_100 | 589382  |

[4545 rows x 3 columns]

By specifying the `value_vars` argument, we can only melt the columns we want.

```
data_long = pd.melt(data,
                    id_vars=["YEAR"],
                    value_vars=["POP_1", "POP_2", "POP_3"],
                    var_name="age",
                    value_name="pop")

data_long
```

|     | YEAR | age   | pop     |
|-----|------|-------|---------|
| 0   | 2016 | POP_1 | 3995008 |
| 1   | 2017 | POP_1 | 3982964 |
| 2   | 2018 | POP_1 | 4068172 |
| 3   | 2019 | POP_1 | 4089881 |
| 4   | 2020 | POP_1 | 4110117 |
| ..  | ...  | ...   | ...     |
| 130 | 2056 | POP_3 | 4401231 |
| 131 | 2057 | POP_3 | 4411893 |
| 132 | 2058 | POP_3 | 4421774 |
| 133 | 2059 | POP_3 | 4430923 |
| 134 | 2060 | POP_3 | 4439404 |

[135 rows x 3 columns]

We can use `str.split()` to remove the prefix `POP_` from the `age` column.

```
data_long["age"] = data_long["age"].apply(lambda x: x.split("_")[1])
data_long
```

```

      YEAR age      pop
0    2016   1  3995008
1    2017   1  3982964
2    2018   1  4068172
3    2019   1  4089881
4    2020   1  4110117
..     ...  ..      ...
130  2056   3  4401231
131  2057   3  4411893
132  2058   3  4421774
133  2059   3  4430923
134  2060   3  4439404

```

```
[135 rows x 3 columns]
```

## 1.2 Billboard top hits for 2000

This dataset contains the song information of the top hits in 2000 and their rank by week. The ranks were organized in multiple columns (e.g., x1st.week, x2nd.week), where each column represents the rank of the song in a week.

```
data = pd.read_csv("tidy_2_bboard.csv")
data
```

```

      year      artist.inverted      track
time \
0    2000  Destiny's Child  Independent Women Part I
03:38
1    2000          Santana      Maria, Maria
04:18
2    2000  Savage Garden  I Knew I Loved You
04:07
3    2000          Madonna      Music
03:45
4    2000  Aguilera, Christina  Come On Over Baby (All I Want Is You)
03:38
..     ...      ...      ...
...
312  2000  Ghostface Killah  Cherchez LaGhost
03:04
313  2000      Smith, Will  Freakin' It
03:58
314  2000  Zombie Nation  Kernkraft 400
03:30

```

```

315 2000 Eastsidaz, The Got Beef
03:58
316 2000 Fragma Toca's Miracle
03:22

  genre date.entered date.peaked x1st.week x2nd.week x3rd.week ...
\
0  Rock      9/23/00    11/18/00        78      63.0      49.0 ...
1  Rock      2/12/00      4/8/00        15       8.0       6.0 ...
2  Rock     10/23/99     1/29/00        71      48.0      43.0 ...
3  Rock      8/12/00     9/16/00        41      23.0      18.0 ...
4  Rock      8/5/00     10/14/00       57      47.0      45.0 ...
..    ...          ...          ...    ...    ...    ...
312  R&B      8/5/00     8/5/00        98      NaN      NaN ...
313  Rap      2/12/00     2/12/00       99      99.0     99.0 ...
314  Rock      9/2/00     9/2/00       99      99.0      NaN ...
315  Rap      7/1/00     7/1/00       99      99.0      NaN ...
316  R&B     10/28/00    10/28/00       99      NaN      NaN ...

  x73rd.week x74th.week x75th.week x76th.week
0          NaN          NaN          NaN          NaN
1          NaN          NaN          NaN          NaN
2          NaN          NaN          NaN          NaN
3          NaN          NaN          NaN          NaN
4          NaN          NaN          NaN          NaN
..    ...          ...          ...          ...
312          NaN          NaN          NaN          NaN
313          NaN          NaN          NaN          NaN
314          NaN          NaN          NaN          NaN
315          NaN          NaN          NaN          NaN
316          NaN          NaN          NaN          NaN

```

[317 rows x 83 columns]

Again, let's identify the `id_vars` and `value_vars`. The `id_vars` are `year`, `artist.inverted`, `track`, `time`, `genre`, `date.entered`, and `date.peaked`. The `value_vars` are the remaining columns.

```

id_vars = data.columns[:7]
data_long = data.melt(id_vars=id_vars, var_name="week", value_name="rank")
data_long = data_long.dropna()
data_long

```

```

  year      artist.inverted      track \
0  2000  Destiny's Child  Independent Women Part I
1  2000      Santana      Maria, Maria
2  2000  Savage Garden  I Knew I Loved You
3  2000      Madonna      Music

```

|       |      |                     |                                       |
|-------|------|---------------------|---------------------------------------|
| 4     | 2000 | Aguilera, Christina | Come On Over Baby (All I Want Is You) |
| ...   | ...  | ...                 | ...                                   |
| 19663 | 2000 | Lonestar            | Amazed                                |
| 19700 | 2000 | Creed               | Higher                                |
| 19980 | 2000 | Lonestar            | Amazed                                |
| 20017 | 2000 | Creed               | Higher                                |
| 20334 | 2000 | Creed               | Higher                                |

|       | time  | genre   | date.entered | date.peaked | week       | rank |
|-------|-------|---------|--------------|-------------|------------|------|
| 0     | 03:38 | Rock    | 9/23/00      | 11/18/00    | x1st.week  | 78.0 |
| 1     | 04:18 | Rock    | 2/12/00      | 4/8/00      | x1st.week  | 15.0 |
| 2     | 04:07 | Rock    | 10/23/99     | 1/29/00     | x1st.week  | 71.0 |
| 3     | 03:45 | Rock    | 8/12/00      | 9/16/00     | x1st.week  | 41.0 |
| 4     | 03:38 | Rock    | 8/5/00       | 10/14/00    | x1st.week  | 57.0 |
| ...   | ...   | ...     | ...          | ...         | ...        | ...  |
| 19663 | 04:25 | Country | 6/5/99       | 3/4/00      | x63rd.week | 45.0 |
| 19700 | 05:16 | Rock    | 9/11/99      | 7/22/00     | x63rd.week | 50.0 |
| 19980 | 04:25 | Country | 6/5/99       | 3/4/00      | x64th.week | 50.0 |
| 20017 | 05:16 | Rock    | 9/11/99      | 7/22/00     | x64th.week | 50.0 |
| 20334 | 05:16 | Rock    | 9/11/99      | 7/22/00     | x65th.week | 49.0 |

[5307 rows x 9 columns]

Refine the **week** column by removing the prefix **x** and the suffix **.week**.

```
data_long["week"] = data_long["week"].apply(lambda x: re.findall(r"\d+",
x)[0]).astype(int)
data_long["rank"] = data_long["rank"].astype(int)
data_long
```

|       | year | artist.inverted     | track \                               |
|-------|------|---------------------|---------------------------------------|
| 0     | 2000 | Destiny's Child     | Independent Women Part I              |
| 1     | 2000 | Santana             | Maria, Maria                          |
| 2     | 2000 | Savage Garden       | I Knew I Loved You                    |
| 3     | 2000 | Madonna             | Music                                 |
| 4     | 2000 | Aguilera, Christina | Come On Over Baby (All I Want Is You) |
| ...   | ...  | ...                 | ...                                   |
| 19663 | 2000 | Lonestar            | Amazed                                |
| 19700 | 2000 | Creed               | Higher                                |
| 19980 | 2000 | Lonestar            | Amazed                                |
| 20017 | 2000 | Creed               | Higher                                |
| 20334 | 2000 | Creed               | Higher                                |

|   | time  | genre | date.entered | date.peaked | week | rank |
|---|-------|-------|--------------|-------------|------|------|
| 0 | 03:38 | Rock  | 9/23/00      | 11/18/00    | 1    | 78   |
| 1 | 04:18 | Rock  | 2/12/00      | 4/8/00      | 1    | 15   |
| 2 | 04:07 | Rock  | 10/23/99     | 1/29/00     | 1    | 71   |

```

3      03:45      Rock      8/12/00      9/16/00      1      41
4      03:38      Rock      8/5/00       10/14/00     1      57
...      ...      ...      ...      ...      ...      ...
19663  04:25  Country      6/5/99       3/4/00     63      45
19700  05:16      Rock      9/11/99     7/22/00     63      50
19980  04:25  Country      6/5/99     3/4/00     64      50
20017  05:16      Rock      9/11/99     7/22/00     64      50
20334  05:16      Rock      9/11/99     7/22/00     65      49

```

```
[5307 rows x 9 columns]
```

## 2. Multiple variables are stored in one column

### 2.1 Tuberculosis (TB) dataset

This is a dataset describing the number of tuberculosis cases in different countries (column **country**) and years (column **year**). The remaining columns were coded in a format of **new\_sp\_xyyyy** where **x** is the gender and **yyyy** is numeric age range. For example, **new\_sp\_m014** kept the TB cases of male patients in the age range of 0-14.

```
data = pd.read_csv("tidy_3_tb.csv")
data
```

```

      country  year  new_sp  new_sp_m04  new_sp_m14  new_sp_m014  \
0          AD  1989     NaN          NaN          NaN          NaN
1          AD  1990     NaN          NaN          NaN          NaN
2          AD  1991     NaN          NaN          NaN          NaN
3          AD  1992     NaN          NaN          NaN          NaN
4          AD  1993    15.0          NaN          NaN          NaN
...      ...    ...      ...      ...      ...      ...
5764       ZW  2004  14581.0          NaN          NaN        187.0
5765       ZW  2005  13155.0          NaN          NaN        210.0
5766       ZW  2006  12718.0          NaN          NaN        215.0
5767       ZW  2007  10583.0         6.0        132.0        138.0
5768       ZW  2008   9830.0          NaN          NaN        127.0

      new_sp_m1524  new_sp_m2534  new_sp_m3544  new_sp_m4554  ...
new_sp_f04  \
0          NaN          NaN          NaN          NaN  ...
NaN
1          NaN          NaN          NaN          NaN  ...
NaN
2          NaN          NaN          NaN          NaN  ...
NaN
3          NaN          NaN          NaN          NaN  ...
NaN
4          NaN          NaN          NaN          NaN  ...

```

```

NaN
...      ...      ...      ...      ...      ...
...
5764      833.0      2908.0      2298.0      1056.0      ...
NaN
5765      837.0      2264.0      1855.0      762.0      ...
NaN
5766      736.0      2391.0      1939.0      896.0      ...
NaN
5767      500.0      3693.0      0.0      716.0      ...
7.0
5768      614.0      0.0      3316.0      704.0      ...
NaN

      new_sp_f514  new_sp_f014  new_sp_f1524  new_sp_f2534  new_sp_f3544
\
0      NaN      NaN      NaN      NaN      NaN
1      NaN      NaN      NaN      NaN      NaN
2      NaN      NaN      NaN      NaN      NaN
3      NaN      NaN      NaN      NaN      NaN
4      NaN      NaN      NaN      NaN      NaN
...      ...      ...      ...      ...      ...
5764      NaN      225.0      1140.0      2858.0      1565.0
5765      NaN      269.0      1136.0      2242.0      1255.0
5766      NaN      237.0      1020.0      2424.0      1355.0
5767      178.0      185.0      739.0      3311.0      0.0
5768      NaN      145.0      840.0      0.0      2890.0

      new_sp_f4554  new_sp_f5564  new_sp_f65  new_sp_fu
0      NaN      NaN      NaN      NaN
1      NaN      NaN      NaN      NaN
2      NaN      NaN      NaN      NaN
3      NaN      NaN      NaN      NaN
4      NaN      NaN      NaN      NaN
...      ...      ...      ...      ...
5764      622.0      214.0      111.0      NaN
5765      578.0      193.0      603.0      NaN
5766      632.0      230.0      96.0      NaN
5767      553.0      213.0      90.0      NaN
5768      467.0      174.0      105.0      0.0

[5769 rows x 23 columns]

```

Let's first tidy the columns of TB cases.

```

data_long = data.melt(id_vars=["country", "year"]).dropna()
data_long

```



|        | country | year | variable  | value |
|--------|---------|------|-----------|-------|
| 4      | AD      | 1993 | new_sp    | 15.0  |
| 5      | AD      | 1994 | new_sp    | 24.0  |
| 6      | AD      | 1996 | new_sp    | 8.0   |
| 7      | AD      | 1997 | new_sp    | 17.0  |
| 8      | AD      | 1998 | new_sp    | 1.0   |
| ...    | ...     | ...  | ...       | ...   |
| 120964 | VU      | 2008 | new_sp_fu | 0.0   |
| 121038 | YE      | 2008 | new_sp_fu | 0.0   |
| 121092 | ZA      | 2008 | new_sp_fu | 0.0   |
| 121119 | ZM      | 2008 | new_sp_fu | 0.0   |
| 121148 | ZW      | 2008 | new_sp_fu | 0.0   |

[38619 rows x 4 columns]

Remove redundant columns.

```
data_long2 = data_long.query("variable not in ['new_sp', 'new_sp_mu',
'new_sp_fu']")
data_long2
```

|        | country | year | variable   | value |
|--------|---------|------|------------|-------|
| 5784   | AD      | 2005 | new_sp_m04 | 0.0   |
| 5785   | AD      | 2006 | new_sp_m04 | 0.0   |
| 5787   | AD      | 2008 | new_sp_m04 | 0.0   |
| 5811   | AE      | 2006 | new_sp_m04 | 0.0   |
| 5812   | AE      | 2007 | new_sp_m04 | 0.0   |
| ...    | ...     | ...  | ...        | ...   |
| 115375 | ZW      | 2004 | new_sp_f65 | 111.0 |
| 115376 | ZW      | 2005 | new_sp_f65 | 603.0 |
| 115377 | ZW      | 2006 | new_sp_f65 | 96.0  |
| 115378 | ZW      | 2007 | new_sp_f65 | 90.0  |
| 115379 | ZW      | 2008 | new_sp_f65 | 105.0 |

[35009 rows x 4 columns]

Then, format the `variable` column to `gender` and `age`.

```
data_long2["age"] = data_long2["variable"].\
    apply(lambda x: re.findall(r"\d+", x.replace("new_sp_", ""))[0])
data_long2["gender"] = data_long2["variable"].\
    apply(lambda x: re.findall(r"[mf]+", x.replace("new_sp_", ""))[0])
data_long2
```

```

      country  year  variable  value  age  gender
5784      AD  2005  new_sp_m04    0.0   04      m
5785      AD  2006  new_sp_m04    0.0   04      m
5787      AD  2008  new_sp_m04    0.0   04      m
5811      AE  2006  new_sp_m04    0.0   04      m
5812      AE  2007  new_sp_m04    0.0   04      m
...
115375     ZW  2004  new_sp_f65  111.0   65      f
115376     ZW  2005  new_sp_f65  603.0   65      f
115377     ZW  2006  new_sp_f65   96.0   65      f
115378     ZW  2007  new_sp_f65   90.0   65      f
115379     ZW  2008  new_sp_f65  105.0   65      f

```

[35009 rows x 6 columns]

### 3. Variables are stored in both rows and columns

This is a scenario where variables are stored in both rows and columns. We will work on the **weather** dataset to demonstrate how to tidy this type of data. Since this dataset did not have a column header, we will use the **columns** attribute to specify the column names.

```

data = pd.read_csv("tidy_4_weather.csv", header=None)
nrow, ncol = data.shape
data.columns = ["idx"] + [i for i in range(1, ncol)]
data

```

```

      idx    1    2    3    4    5    6    7    8    9
...  \
0  MX000017004195504TMIN  150  150  160  150  160  160  160  160  160
...
1  MX000017004195504PRCP    0    0    0    0    0    0    0    0    0
...
2  MX000017004195505TMAX  310  310  310  300  300  300  310  310  310
...
3  MX000017004195505TMIN  200  160  160  150  150  150  160  160  170
...
4  MX000017004195505PRCP    0    0    0    0    0    0    0    0    0
...
..      ...    ...    ...    ...    ...    ...    ...    ...    ...
...
994 MX000017004199004TMIN    0  147  147  150  136    0  157    0    0
...
995 MX000017004199004PRCP    0    0    0    0   23    0    0    0    0
...
996 MX000017004199005TMAX    0  350  362  348  337    0  240    0    0
...

```

|                         |                       |     |     |     |     |     |     |     |     |     |     |
|-------------------------|-----------------------|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 997                     | MX000017004199005TMIN | 168 | 168 | 167 | 167 | 170 | 0   | 132 | 132 | 0   |     |
| ...                     |                       |     |     |     |     |     |     |     |     |     |     |
| 998                     | MX000017004199005PRCP | 0   | 0   | 0   | 0   | 61  | 0   | 254 | 20  | 0   |     |
| ...                     |                       |     |     |     |     |     |     |     |     |     |     |
|                         |                       | 22  | 23  | 24  | 25  | 26  | 27  | 28  | 29  | 30  | 31  |
| 0                       |                       | 170 | 170 | 170 | 180 | 190 | 190 | 170 | 180 | 160 | 0   |
| 1                       |                       | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 6   | 0   |
| 2                       |                       | 330 | 340 | 350 | 330 | 310 | 310 | 320 | 310 | 300 | 290 |
| 3                       |                       | 170 | 190 | 190 | 190 | 180 | 160 | 150 | 170 | 150 | 160 |
| 4                       |                       | 0   | 0   | 0   | 0   | 0   | 142 | 0   | 54  | 0   | 46  |
| ..                      |                       | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 994                     |                       | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 157 | 0   |
| 995                     |                       | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   |
| 996                     |                       | 0   | 297 | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 336 |
| 997                     |                       | 144 | 136 | 136 | 155 | 155 | 0   | 0   | 166 | 179 | 169 |
| 998                     |                       | 0   | 89  | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 3   |
| [999 rows x 32 columns] |                       |     |     |     |     |     |     |     |     |     |     |

You should notice that the column `idx` contains complicated information: `station ID`, `year`, `month`, and `day`. We can use simple indexing to extract the information we want.

```
data["site"] = data["idx"].apply(lambda x: x[:11])
data["year"] = data["idx"].apply(lambda x: x[-10:-6]).astype(int)
data["month"] = data["idx"].apply(lambda x: x[-6:-4]).astype(int)
data["variable"] = data["idx"].apply(lambda x: x[-4:])
data = data.drop("idx", axis=1)
data
```

|     |     |     |     |             |      |      |      |       |     |          |     |     |     |     |
|-----|-----|-----|-----|-------------|------|------|------|-------|-----|----------|-----|-----|-----|-----|
|     | 1   | 2   | 3   | 4           | 5    | 6    | 7    | 8     | 9   | 10       | ... | 26  | 27  | 28  |
| \   |     |     |     |             |      |      |      |       |     |          |     |     |     |     |
| 0   | 150 | 150 | 160 | 150         | 160  | 160  | 160  | 160   | 160 | 170      | ... | 190 | 190 | 170 |
| 1   | 0   | 0   | 0   | 0           | 0    | 0    | 0    | 0     | 0   | 0        | ... | 0   | 0   | 0   |
| 2   | 310 | 310 | 310 | 300         | 300  | 300  | 310  | 310   | 310 | 300      | ... | 310 | 310 | 320 |
| 3   | 200 | 160 | 160 | 150         | 150  | 150  | 160  | 160   | 170 | 170      | ... | 180 | 160 | 150 |
| 4   | 0   | 0   | 0   | 0           | 0    | 0    | 0    | 0     | 0   | 0        | ... | 0   | 142 | 0   |
| ..  | ... | ... | ... | ...         | ...  | ...  | ...  | ...   | ... | ...      | ... | ... | ... | ... |
| 994 | 0   | 147 | 147 | 150         | 136  | 0    | 157  | 0     | 0   | 157      | ... | 0   | 0   | 0   |
| 995 | 0   | 0   | 0   | 0           | 23   | 0    | 0    | 0     | 0   | 0        | ... | 0   | 0   | 0   |
| 996 | 0   | 350 | 362 | 348         | 337  | 0    | 240  | 0     | 0   | 0        | ... | 0   | 0   | 0   |
| 997 | 168 | 168 | 167 | 167         | 170  | 0    | 132  | 132   | 0   | 0        | ... | 155 | 0   | 0   |
| 998 | 0   | 0   | 0   | 0           | 61   | 0    | 254  | 20    | 0   | 5        | ... | 0   | 0   | 0   |
|     | 29  | 30  | 31  |             |      | site | year | month |     | variable |     |     |     |     |
| 0   | 180 | 160 | 0   | MX000017004 | 1955 |      |      | 4     |     | TMIN     |     |     |     |     |
| 1   | 0   | 6   | 0   | MX000017004 | 1955 |      |      | 4     |     | PRCP     |     |     |     |     |

|     |     |     |     |             |      |     |      |
|-----|-----|-----|-----|-------------|------|-----|------|
| 2   | 310 | 300 | 290 | MX000017004 | 1955 | 5   | TMAX |
| 3   | 170 | 150 | 160 | MX000017004 | 1955 | 5   | TMIN |
| 4   | 54  | 0   | 46  | MX000017004 | 1955 | 5   | PRCP |
| ... | ... | ... | ... | ...         | ...  | ... | ...  |
| 994 | 0   | 157 | 0   | MX000017004 | 1990 | 4   | TMIN |
| 995 | 0   | 0   | 0   | MX000017004 | 1990 | 4   | PRCP |
| 996 | 0   | 0   | 336 | MX000017004 | 1990 | 5   | TMAX |
| 997 | 166 | 179 | 169 | MX000017004 | 1990 | 5   | TMIN |
| 998 | 0   | 0   | 3   | MX000017004 | 1990 | 5   | PRCP |

[999 rows x 35 columns]

Now, let's tidy the date information.

```
data2 = data.melt(id_vars=["site", "year", "month", "variable"],
var_name="day", value_name="value")
data2["day"] = data2["day"].astype(int)
data2
```

|       | site        | year | month | variable | day | value |
|-------|-------------|------|-------|----------|-----|-------|
| 0     | MX000017004 | 1955 | 4     | TMIN     | 1   | 150   |
| 1     | MX000017004 | 1955 | 4     | PRCP     | 1   | 0     |
| 2     | MX000017004 | 1955 | 5     | TMAX     | 1   | 310   |
| 3     | MX000017004 | 1955 | 5     | TMIN     | 1   | 200   |
| 4     | MX000017004 | 1955 | 5     | PRCP     | 1   | 0     |
| ...   | ...         | ...  | ...   | ...      | ... | ...   |
| 30964 | MX000017004 | 1990 | 4     | TMIN     | 31  | 0     |
| 30965 | MX000017004 | 1990 | 4     | PRCP     | 31  | 0     |
| 30966 | MX000017004 | 1990 | 5     | TMAX     | 31  | 336   |
| 30967 | MX000017004 | 1990 | 5     | TMIN     | 31  | 169   |
| 30968 | MX000017004 | 1990 | 5     | PRCP     | 31  | 3     |

[30969 rows x 6 columns]

Now, since there are three variables: `tmax`, `tmin`, and `prcp` in the `variable` column, we need to pivot it to the column names.

We can use `df.pivot()` to do this. The `index` argument specifies the columns that should not be pivoted. The `columns` argument specifies the columns that should be pivoted. The `values` argument specifies the column that should be used as the values of the pivoted columns.

```
data3 = data2.pivot(index = ["site", "year", "month", "day"], columns =
"variable", values="value")
data3 = data3.reset_index().dropna()
data3
```

```

variable      site year month day PRCP  TMAX  TMIN
31      MX000017004 1955     5   1   0.0  310.0 200.0
32      MX000017004 1955     5   2   0.0  310.0 160.0
33      MX000017004 1955     5   3   0.0  310.0 160.0
34      MX000017004 1955     5   4   0.0  300.0 150.0
35      MX000017004 1955     5   5   0.0  300.0 150.0
...      ...      ...      ...   ...   ...   ...
10411     MX000017004 1990     5  27   0.0   0.0   0.0
10412     MX000017004 1990     5  28   0.0   0.0   0.0
10413     MX000017004 1990     5  29   0.0   0.0  166.0
10414     MX000017004 1990     5  30   0.0   0.0  179.0
10415     MX000017004 1990     5  31   3.0  336.0 169.0

```

[10230 rows x 7 columns]

## 4. Multiple types of observational units are stored in the same table

Let's revisit the billboard dataset that we worked in the second type of messy data. Although we have tidied the dataset, we can still see many repeated values in the data. For example, if we only check records of the song "Maria, Maria", we can see that multiple columns, such as `artist.inverted` and `track`, were repeated for each week.

```

data = pd.read_csv("tidy_5_bboard_long.csv")
data

```

```

      year  artist.inverted  track \
0      2000  Destiny's Child  Independent Women Part I
1      2000           Santana  Maria, Maria
2      2000  Savage Garden  I Knew I Loved You
3      2000           Madonna  Music
4      2000  Aguilera, Christina  Come On Over Baby (All I Want Is You)
...      ...      ...      ...
24087  2000  Ghostface Killah  Cherchez LaGhost
24088  2000      Smith, Will  Freakin' It
24089  2000  Zombie Nation  Kernkraft 400
24090  2000  Eastsidaz, The  Got Beef
24091  2000           Fragma  Toca's Miracle

      time genre date.entered date.peaked  week  rank
0      03:38  Rock      9/23/00    11/18/00     1  78.0
1      04:18  Rock      2/12/00      4/8/00     1  15.0
2      04:07  Rock     10/23/99     1/29/00     1  71.0
3      03:45  Rock      8/12/00     9/16/00     1  41.0
4      03:38  Rock      8/5/00    10/14/00     1  57.0

```

```

...      ...      ...      ...      ...      ...      ...
24087  03:04  R&B      8/5/00      8/5/00      76      NaN
24088  03:58  Rap       2/12/00     2/12/00     76      NaN
24089  03:30  Rock      9/2/00      9/2/00     76      NaN
24090  03:58  Rap       7/1/00      7/1/00     76      NaN
24091  03:22  R&B     10/28/00    10/28/00    76      NaN

```

```
[24092 rows x 9 columns]
```

```
data.query("track == 'Maria, Maria'")
```

```

      year artist.inverted      track      time genre date.entered \
1      2000      Santana  Maria, Maria  04:18  Rock      2/12/00
318    2000      Santana  Maria, Maria  04:18  Rock      2/12/00
635    2000      Santana  Maria, Maria  04:18  Rock      2/12/00
952    2000      Santana  Maria, Maria  04:18  Rock      2/12/00
1269   2000      Santana  Maria, Maria  04:18  Rock      2/12/00
...      ...      ...      ...      ...      ...      ...
22508  2000      Santana  Maria, Maria  04:18  Rock      2/12/00
22825  2000      Santana  Maria, Maria  04:18  Rock      2/12/00
23142  2000      Santana  Maria, Maria  04:18  Rock      2/12/00
23459  2000      Santana  Maria, Maria  04:18  Rock      2/12/00
23776  2000      Santana  Maria, Maria  04:18  Rock      2/12/00

```

```

      date.peakd  week  rank
1      4/8/00      1  15.0
318    4/8/00      2   8.0
635    4/8/00      3   6.0
952    4/8/00      4   5.0
1269   4/8/00      5   2.0
...      ...      ...      ...
22508    4/8/00     72  NaN
22825    4/8/00     73  NaN
23142    4/8/00     74  NaN
23459    4/8/00     75  NaN
23776    4/8/00     76  NaN

```

```
[76 rows x 9 columns]
```

You can observe duplicated information being stored for the song "Maria, Maria". We need a separate table to store the song information only once.

### Get the columns that contain only song information

```
header_idx = ["year", "artist.inverted", "track", "time", "genre",
              "date.entered", "date.peaked"]
data_song = data.loc[:, header_idx]
data_song
```

```

      year      artist.inverted      track \
0      2000      Destiny's Child      Independent Women Part I
1      2000              Santana      Maria, Maria
2      2000      Savage Garden      I Knew I Loved You
3      2000              Madonna      Music
4      2000  Aguilera, Christina  Come On Over Baby (All I Want Is You)
...      ...              ...      ...
24087  2000      Ghostface Killah      Cherchez LaGhost
24088  2000          Smith, Will      Freakin' It
24089  2000      Zombie Nation      Kernkraft 400
24090  2000      Eastsidaz, The      Got Beef
24091  2000              Fragma      Toca's Miracle

      time genre date.entered date.peaked
0      03:38  Rock      9/23/00      11/18/00
1      04:18  Rock      2/12/00       4/8/00
2      04:07  Rock     10/23/99      1/29/00
3      03:45  Rock      8/12/00      9/16/00
4      03:38  Rock      8/5/00      10/14/00
...      ...      ...      ...      ...
24087  03:04  R&B      8/5/00      8/5/00
24088  03:58  Rap      2/12/00      2/12/00
24089  03:30  Rock      9/2/00      9/2/00
24090  03:58  Rap      7/1/00      7/1/00
24091  03:22  R&B     10/28/00     10/28/00

[24092 rows x 7 columns]
```

Drop the duplicated rows using `pd.drop_duplicates()`. Check the row count of the new dataframe.

```
data_song = data_song.drop_duplicates()
data_song
```

```

      year      artist.inverted      track
time \
0      2000      Destiny's Child      Independent Women Part I
03:38
1      2000              Santana      Maria, Maria
04:18
```

```

2    2000    Savage Garden    I Knew I Loved You
04:07
3    2000    Madonna    Music
03:45
4    2000    Aguilera, Christina    Come On Over Baby (All I Want Is You)
03:38
..    ...    ...
...
312  2000    Ghostface Killah    Cherchez LaGhost
03:04
313  2000    Smith, Will    Freakin' It
03:58
314  2000    Zombie Nation    Kernkraft 400
03:30
315  2000    Eastsidaz, The    Got Beef
03:58
316  2000    Fragma    Toca's Miracle
03:22

```

```

    genre date.entered date.peaked
0    Rock    9/23/00    11/18/00
1    Rock    2/12/00    4/8/00
2    Rock    10/23/99    1/29/00
3    Rock    8/12/00    9/16/00
4    Rock    8/5/00    10/14/00
..    ...    ...
312  R&B    8/5/00    8/5/00
313  Rap    2/12/00    2/12/00
314  Rock    9/2/00    9/2/00
315  Rap    7/1/00    7/1/00
316  R&B    10/28/00    10/28/00

```

```
[317 rows x 7 columns]
```

We need to create a numeric index for the song table, because we will use it as a **foreign key** in the ranking table. Use `pd.reset_index()` to create a new column **index** and use it as the index of the song table.

```

data_song = data_song.reset_index()
data_song

```

```

    index  year    artist.inverted
track \
0      0  2000    Destiny's Child    Independent Women Part
I
1      1  2000    Santana    Maria,
Maria

```



```

2          2  2000          Savage Garden          I Knew I Loved
You
3          3  2000          Madonna
Music
4          4  2000  Aguilera, Christina  Come On Over Baby (All I Want Is
You)
..        ...    ...
...
312       312  2000          Ghostface Killah          Cherchez
LaGhost
313       313  2000          Smith, Will          Freakin'
It
314       314  2000          Zombie Nation          Kernkraft
400
315       315  2000          Eastsidaz, The          Got
Beef
316       316  2000          Fragma          Toca's
Miracle

```

```

      time genre date.entered date.peaked
0    03:38  Rock      9/23/00    11/18/00
1    04:18  Rock      2/12/00      4/8/00
2    04:07  Rock    10/23/99     1/29/00
3    03:45  Rock      8/12/00     9/16/00
4    03:38  Rock      8/5/00    10/14/00
..     ...    ...
312   03:04  R&B      8/5/00     8/5/00
313   03:58  Rap      2/12/00     2/12/00
314   03:30  Rock      9/2/00     9/2/00
315   03:58  Rap      7/1/00     7/1/00
316   03:22  R&B    10/28/00    10/28/00

```

```
[317 rows x 8 columns]
```

Merge the `data_song` and the original `data` using `pd.merge()`. Now, we get the numeric index for each song in the ranking table.

<https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.DataFrame.merge.html>

```

data_merge = pd.merge(data_song, data, on=header_idx)
data_merge

```

```

      index  year  artist.inverted      track  time genre
\
0         0  2000  Destiny's Child  Independent Women Part I  03:38  Rock
1         0  2000  Destiny's Child  Independent Women Part I  03:38  Rock
2         0  2000  Destiny's Child  Independent Women Part I  03:38  Rock

```

```

3          0  2000  Destiny's Child  Independent Women Part I  03:38  Rock
4          0  2000  Destiny's Child  Independent Women Part I  03:38  Rock
...
24087     316  2000          Fragma          Toca's Miracle  03:22  R&B
24088     316  2000          Fragma          Toca's Miracle  03:22  R&B
24089     316  2000          Fragma          Toca's Miracle  03:22  R&B
24090     316  2000          Fragma          Toca's Miracle  03:22  R&B
24091     316  2000          Fragma          Toca's Miracle  03:22  R&B

```

```

      date.entered date.peakd  week  rank
0      9/23/00    11/18/00     1  78.0
1      9/23/00    11/18/00     2  63.0
2      9/23/00    11/18/00     3  49.0
3      9/23/00    11/18/00     4  33.0
4      9/23/00    11/18/00     5  23.0
...
24087    10/28/00    10/28/00    72   NaN
24088    10/28/00    10/28/00    73   NaN
24089    10/28/00    10/28/00    74   NaN
24090    10/28/00    10/28/00    75   NaN
24091    10/28/00    10/28/00    76   NaN

```

[24092 rows x 10 columns]

### Get the columns that contain only song information

Finally, we can drop the song information to get the ranking table, which is indexed by the numeric index column: `index`.

```

data_rank = data_merge.loc[:, ["index", "week", "rank"]]
data_rank

```

```

      index  week  rank
0         0     1  78.0
1         0     2  63.0
2         0     3  49.0
3         0     4  33.0
4         0     5  23.0
...
24087    316    72   NaN
24088    316    73   NaN
24089    316    74   NaN
24090    316    75   NaN
24091    316    76   NaN

```

[24092 rows x 3 columns]

We can compare the number of values we need before and after the transformation.

```
size_original = data.size
size_new = data_rank.size + data_song.size

print("The number of elements (number of rows times number of columns) in
data is ", size_original)
print("New data has ", size_new, "elements")
print("Compression ratio is ", size_new / size_original)
```

```
The number of elements (number of rows times number of columns) in data is
216828
New data has 74812 elements
Compression ratio is 0.34502923976608185
```

## 5. A single observational unit is stored in multiple tables

This is a case where a single observational unit is stored in multiple tables organized by different variables. In this example, we will work on the **babynames** dataset. This dataset consists of multiple files, each of which contains the baby names and their proportions. The files are organized by year and gender in a format of **babynames\_YYYY\_gender.csv**, where **YYYY** is the year and **gender** is the gender (i.e., boy or girl).

Let's observe the file naming pattern first

```
os.listdir("tidy_6_babynames")
```

```
['babynames_1887_boy.csv',
 'babynames_1897_boy.csv',
 'babynames_1959_girl.csv',
 'babynames_1958_girl.csv',
 'babynames_1946_boy.csv',
 'babynames_1956_boy.csv',
 'babynames_1982_girl.csv',
 ...,
 'babynames_1931_boy.csv',
 'babynames_1921_boy.csv',
 'babynames_1974_girl.csv',
 'babynames_1975_girl.csv',
 'babynames_1988_boy.csv',
 'babynames_1998_boy.csv']
```

We can try to open one of them to inspect the data.

```
pd.read_csv(os.path.join("tidy_6_babynames", "babynames_1887_boy.csv"))
```

|     | name    | percent  |
|-----|---------|----------|
| 0   | John    | 0.074181 |
| 1   | William | 0.068344 |
| 2   | James   | 0.043617 |
| 3   | George  | 0.039190 |
| 4   | Charles | 0.036875 |
| ..  | ...     | ...      |
| 995 | Jessee  | 0.000046 |
| 996 | Jewel   | 0.000046 |
| 997 | Jodie   | 0.000046 |
| 998 | Lars    | 0.000046 |
| 999 | Laurel  | 0.000046 |

[1000 rows x 2 columns]

You will find that there is no year or gender information in the file, we need to extract them from the file name. Let's start with a single file to experiment the extraction process.

```
filename = "babynames_1887_boy.csv"
year = re.findall(r"\d+", filename)[0]
gender = re.findall(r"[a-z]+\.", filename)[0].replace(".", "")
data = pd.read_csv(os.path.join("tidy_6_babynames", filename))
data["year"] = year
data["gender"] = gender
data
```

|     | name    | percent  | year | gender |
|-----|---------|----------|------|--------|
| 0   | John    | 0.074181 | 1887 | boy    |
| 1   | William | 0.068344 | 1887 | boy    |
| 2   | James   | 0.043617 | 1887 | boy    |
| 3   | George  | 0.039190 | 1887 | boy    |
| 4   | Charles | 0.036875 | 1887 | boy    |
| ..  | ...     | ...      | ...  | ...    |
| 995 | Jessee  | 0.000046 | 1887 | boy    |
| 996 | Jewel   | 0.000046 | 1887 | boy    |
| 997 | Jodie   | 0.000046 | 1887 | boy    |
| 998 | Lars    | 0.000046 | 1887 | boy    |
| 999 | Laurel  | 0.000046 | 1887 | boy    |

[1000 rows x 4 columns]

Now, we can iteratively extract needed information from each file and store them into one csv file. We can create an empty file with only the header defined.

```
FILE_OUT = "tidy_6_babynames.csv"
with open(FILE_OUT, "w") as f:
    f.write("name,percent,year,gender\n")
```

Next, let's test the extraction process on the first five files. Try to define constant variables instead of hard-coding the values every time.

```
DIR_DATA = "tidy_6_babynames"
ls_files = os.listdir(DIR_DATA)
for filename in ls_files:
    year = re.findall(r"\d+", filename)[0]
    gender = re.findall(r"[a-z]+\.", filename)[0].replace(".", "")
    data = pd.read_csv(os.path.join(DIR_DATA, filename))
    data["year"] = year
    data["gender"] = gender
    data.to_csv(FILE_OUT, mode="a", header=False, index=False)
```

Examine the extracted data.

```
data = pd.read_csv(FILE_OUT)
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 258000 entries, 0 to 257999
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   name        258000 non-null object
1   percent     258000 non-null float64
2   year        258000 non-null int64
3   gender      258000 non-null object
dtypes: float64(1), int64(1), object(2)
memory usage: 7.9+ MB
```

```
data["year"].value_counts()
```

```

1887    2000
1977    2000
1972    2000
1986    2000
1996    2000
...
1913    2000
1906    2000
1907    2000
2000    2000
1991    2000
Name: year, Length: 129, dtype: int64

```

```
data["gender"].value_counts()
```

```

boy      129000
girl     129000
Name: gender, dtype: int64

```

```
data.groupby(["year", "gender"]).agg({"percent": "sum"})
```

```

      percent
year gender
1880 boy    0.930746
      girl  0.934546
1881 boy    0.930439
      girl  0.932690
1882 boy    0.927532
...      ...
2006 girl   0.684830
2007 boy    0.801105
      girl   0.677453
2008 boy    0.795414
      girl   0.672516

[258 rows x 1 columns]

```

## 6. Case study

We will use a case study to illustrate the advantages of tidying data. The dataset `tidy_X.csv` contains the individual-level mortality from Mexico. The columns include the following:

- `sex`: the gender of the deceased
- `age`: the age of the deceased
- `yod`: the year of death
- `mod`: the month of death
- `dod`: the day of death
- `hod`: the hour of death
- `cod`: the cause of death

The goal is to find causes of death with unusual temporal patterns within a day.

## 6.1 Basic data exploration

Let's read the data first.

```
data = pd.read_csv("tidy_X.csv")
data
```

|        | sex | age | yod  | mod | dod | hod | cod |
|--------|-----|-----|------|-----|-----|-----|-----|
| 0      | 1   | 90  | 2008 | 1   | 7   | 20  | F17 |
| 1      | 1   | 72  | 2008 | 1   | 13  | 14  | I05 |
| 2      | 1   | 49  | 2008 | 1   | 12  | 20  | K65 |
| 3      | 2   | 79  | 2008 | 1   | 20  | 10  | I38 |
| 4      | 1   | 15  | 2008 | 1   | 1   | 15  | N18 |
| ...    | ... | ... | ...  | ... | ... | ... | ... |
| 528323 | 1   | 1   | 2008 | 10  | 6   | 12  | P22 |
| 528324 | 2   | 20  | 2008 | 10  | 18  | 20  | Q24 |
| 528325 | 2   | 3   | 2008 | 11  | 11  | 19  | P22 |
| 528326 | 1   | 24  | 2008 | 9   | 25  | 12  | P22 |
| 528327 | 1   | 2   | 2008 | 9   | 22  | 16  | P26 |

[528328 rows x 7 columns]

Check the data types of the columns.

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 528328 entries, 0 to 528327
Data columns (total 7 columns):
#   Column  Non-Null Count  Dtype
---

```

```

0    sex    528328 non-null int64
1    age    528328 non-null int64
2    yod    528328 non-null int64
3    mod    528328 non-null int64
4    dod    528328 non-null int64
5    hod    528328 non-null int64
6    cod    528328 non-null object
dtypes: int64(6), object(1)
memory usage: 28.2+ MB

```

```
data.describe()
```

```

              sex              age              yod              mod  \
count  528328.000000  528328.000000  528328.000000  528328.000000
mean         1.443461         61.246593         2007.950735         6.490457
std          0.496794         24.611694          2.881578         3.554002
min          1.000000         1.000000          0.000000         0.000000
25%          1.000000         48.000000         2008.000000         3.000000
50%          1.000000         67.000000         2008.000000         6.000000
75%          2.000000         80.000000         2008.000000        10.000000
max          2.000000         99.000000         2008.000000        12.000000

              dod              hod
count  528328.000000  528328.000000
mean         15.738475         11.701500
std          8.826922          6.763691
min          0.000000          0.000000
25%          8.000000          6.000000
50%         16.000000         12.000000
75%         23.000000         17.000000
max         31.000000         23.000000

```

What if we want to examine the difference between the genders over the day? We can group `hod` and `sex` to calculate the counts of deaths in each hour.

```

data_grp = data.groupby(["hod", "sex"]).agg(size=("hod", lambda x:
len(x))).reset_index()
data_grp.query("sex == 1")

```

```

   hod  sex  size
0     0    1  12224
2     1    1  11376
4     2    1  10771


```



|    |    |   |       |
|----|----|---|-------|
| 6  | 3  | 1 | 10906 |
| 8  | 4  | 1 | 11140 |
| 10 | 5  | 1 | 11943 |
| 12 | 6  | 1 | 12912 |
| 14 | 7  | 1 | 11843 |
| 16 | 8  | 1 | 11933 |
| 18 | 9  | 1 | 12187 |
| 20 | 10 | 1 | 13281 |
| 22 | 11 | 1 | 12977 |
| 24 | 12 | 1 | 12770 |
| 26 | 13 | 1 | 12647 |
| 28 | 14 | 1 | 12595 |
| 30 | 15 | 1 | 12618 |
| 32 | 16 | 1 | 13050 |
| 34 | 17 | 1 | 13004 |
| 36 | 18 | 1 | 13417 |
| 38 | 19 | 1 | 12743 |
| 40 | 20 | 1 | 12658 |
| 42 | 21 | 1 | 11595 |
| 44 | 22 | 1 | 11413 |
| 46 | 23 | 1 | 12032 |

We can use `plotnine`, a Python version of `ggplot2`, to visualize the data.

```
from plotnine import *
ggplot(data_grp, aes(x="hod", y="size", color="factor(sex)")) +
geom_point() + geom_line()
```



```
<ggplot: (706404852)>
```

## 6.2 Count the number of deaths in each hour

We need to tidy the data by `hod` and `cod` to get the number of deaths in each hour. Before we do that, we can use `df.query()` to check what result we will get. For example, there should be five records of `hod=0`, and `cod=A06`.

```
data.query("hod == 0 and cod == 'A06'")
```

|        | sex | age | yod  | mod | dod | hod | cod |
|--------|-----|-----|------|-----|-----|-----|-----|
| 19640  | 2   | 88  | 2008 | 1   | 10  | 0   | A06 |
| 298031 | 1   | 56  | 2008 | 12  | 28  | 0   | A06 |
| 395562 | 2   | 22  | 2008 | 12  | 20  | 0   | A06 |
| 412156 | 1   | 83  | 2008 | 1   | 5   | 0   | A06 |
| 502497 | 1   | 2   | 2008 | 7   | 19  | 0   | A06 |

Then use `df.groupby()` to group the data by `hod` and `cod` and aggregate the data by `size` to count the number of deaths. We should see "5" in the `freq_by_hodcod` column as we queried above.

```
data_grp = data.groupby(["hod", "cod"]).agg(freq_by_hodcod=("hod",
"size")).reset_index()
```

```
data_grp
```

```

      hod  cod  freq_by_hodcod
0      0  A02      1
1      0  A04      6
2      0  A06      5
3      0  A09     87
4      0  A15      7
...    ...  ...
16171   23  Y34     28
16172   23  Y57      3
16173   23  Y83      7
16174   23  Y86     16
16175   23  Y89      5

```

```
[16176 rows x 3 columns]
```

It is equivalent to use `lambda` function to define the counting function.

```

data_grp = data.groupby(["hod", "cod"]).agg(freq_by_hodcod=("hod", lambda
x: len(x))).reset_index()
data_grp

```

```

      hod  cod  freq_by_hodcod
0      0  A02      1
1      0  A04      6
2      0  A06      5
3      0  A09     91
4      0  A15      7
...    ...  ...
16702   99  Y33      2
16703   99  Y34     99
16704   99  Y40      1
16705   99  Y57      3
16706   99  Y86      2

```

```
[16707 rows x 3 columns]
```

### 6.3 The proportion of deaths in each cause by hour

Now, let's continue with the `data_grp` to compute proportion of deaths of each `cod` given `hod`. We need to determine the denominator of the proportion first. The denominator is the total number of deaths in each `cod` across all `hod`.

Again, we can use `df.query()` to preview what we will get.

```
data_grp.query("cod == 'A03'")
```

|       | hod | cod | freq_by_hodcod |
|-------|-----|-----|----------------|
| 5257  | 8   | A03 | 1              |
| 5922  | 9   | A03 | 1              |
| 6611  | 10  | A03 | 1              |
| 8014  | 12  | A03 | 1              |
| 11503 | 17  | A03 | 1              |
| 12204 | 18  | A03 | 1              |
| 14210 | 21  | A03 | 1              |

Then use `df.groupby()` to group the data by `cod` and aggregate the data by `sum` to get the total number of deaths in each `cod`.

```
sum_cod = data_grp.groupby(["cod"]).agg(sum_by_cod=("freq_by_hodcod",
"sum")).reset_index()
sum_cod
```

|      | cod | sum_by_cod |
|------|-----|------------|
| 0    | A01 | 51         |
| 1    | A02 | 62         |
| 2    | A03 | 7          |
| 3    | A04 | 144        |
| 4    | A05 | 20         |
| ...  | ... | ...        |
| 1192 | Y85 | 4          |
| 1193 | Y86 | 363        |
| 1194 | Y87 | 2          |
| 1195 | Y88 | 5          |
| 1196 | Y89 | 39         |

[1197 rows x 2 columns]

We use `df.merge()` to merge the `data_grp` and the `sum_cod` to put the sum of deaths in each `cod` into the `data_grp` table.

```
data_grp2 = pd.merge(data_grp, sum_cod)
data_grp2
```

|       | hod | cod | freq_by_hodcod | sum_by_cod |
|-------|-----|-----|----------------|------------|
| 0     | 0   | A02 | 1              | 62         |
| 1     | 1   | A02 | 3              | 62         |
| 2     | 2   | A02 | 9              | 62         |
| 3     | 3   | A02 | 1              | 62         |
| 4     | 4   | A02 | 3              | 62         |
| ...   | ... | ... | ...            | ...        |
| 16171 | 22  | Y52 | 2              | 2          |
| 16172 | 23  | D24 | 1              | 1          |
| 16173 | 23  | N88 | 1              | 1          |
| 16174 | 23  | O67 | 1              | 1          |
| 16175 | 23  | V33 | 1              | 1          |

[16176 rows x 4 columns]

Then, with the columns `freq_by_hodcod` as the numerator and `sum_by_cod` as the denominator, we can compute the proportion of deaths in each `hod` given `cod`.

```
data_grp2["prop_by_hodcod"] = data_grp2["freq_by_hodcod"] /
data_grp2["sum_by_cod"]
data_grp2
```

|       | hod | cod | freq_by_hodcod | sum_by_cod | prop_by_hodcod |
|-------|-----|-----|----------------|------------|----------------|
| 0     | 0   | A02 | 1              | 62         | 0.016129       |
| 1     | 1   | A02 | 3              | 62         | 0.048387       |
| 2     | 2   | A02 | 9              | 62         | 0.145161       |
| 3     | 3   | A02 | 1              | 62         | 0.016129       |
| 4     | 4   | A02 | 3              | 62         | 0.048387       |
| ...   | ... | ... | ...            | ...        | ...            |
| 16171 | 22  | Y52 | 2              | 2          | 1.000000       |
| 16172 | 23  | D24 | 1              | 1          | 1.000000       |
| 16173 | 23  | N88 | 1              | 1          | 1.000000       |
| 16174 | 23  | O67 | 1              | 1          | 1.000000       |
| 16175 | 23  | V33 | 1              | 1          | 1.000000       |

[16176 rows x 5 columns]

## 6.4 The proportion of deaths in each hour

Next, to know if a cause of death has unusual temporal patterns, we need to compare the proportion of deaths in each hour with the proportion of deaths in each hour across all causes of death. We can use the `data_grp` to compute the sum of deaths in each hour first.

```
sum_hod = data_grp2.groupby(["hod"]).agg(sum_by_hod=("freq_by_hodcod",
"sum")).reset_index()
sum_hod
```

|    | hod | sum_by_hod |
|----|-----|------------|
| 0  | 0   | 20072      |
| 1  | 1   | 20248      |
| 2  | 2   | 18806      |
| 3  | 3   | 19532      |
| 4  | 4   | 20069      |
| 5  | 5   | 21883      |
| 6  | 6   | 23536      |
| 7  | 7   | 21619      |
| 8  | 8   | 21713      |
| 9  | 9   | 22223      |
| 10 | 10  | 24093      |
| 11 | 11  | 23627      |
| 12 | 12  | 23172      |
| 13 | 13  | 23058      |
| 14 | 14  | 22786      |
| 15 | 15  | 23047      |
| 16 | 16  | 23622      |
| 17 | 17  | 23395      |
| 18 | 18  | 24093      |
| 19 | 19  | 22681      |
| 20 | 20  | 22702      |
| 21 | 21  | 20813      |
| 22 | 22  | 20298      |
| 23 | 23  | 21240      |

Then, we sum the `sum_by_hod` to obtain the total number of deaths as the denominator.

```
sum_hod["sum"] = sum_hod["sum_by_hod"].sum()
sum_hod
```

|   | hod | sum_by_hod | sum    |
|---|-----|------------|--------|
| 0 | 0   | 20072      | 528328 |
| 1 | 1   | 20248      | 528328 |
| 2 | 2   | 18806      | 528328 |
| 3 | 3   | 19532      | 528328 |
| 4 | 4   | 20069      | 528328 |
| 5 | 5   | 21883      | 528328 |
| 6 | 6   | 23536      | 528328 |
| 7 | 7   | 21619      | 528328 |

|    |    |       |        |
|----|----|-------|--------|
| 8  | 8  | 21713 | 528328 |
| 9  | 9  | 22223 | 528328 |
| 10 | 10 | 24093 | 528328 |
| 11 | 11 | 23627 | 528328 |
| 12 | 12 | 23172 | 528328 |
| 13 | 13 | 23058 | 528328 |
| 14 | 14 | 22786 | 528328 |
| 15 | 15 | 23047 | 528328 |
| 16 | 16 | 23622 | 528328 |
| 17 | 17 | 23395 | 528328 |
| 18 | 18 | 24093 | 528328 |
| 19 | 19 | 22681 | 528328 |
| 20 | 20 | 22702 | 528328 |
| 21 | 21 | 20813 | 528328 |
| 22 | 22 | 20298 | 528328 |
| 23 | 23 | 21240 | 528328 |

Finally, we can compute the proportion of deaths in each hour.

```
sum_hod["prop_by_hod"] = sum_hod["sum_by_hod"] / sum_hod["sum"]
sum_hod = sum_hod.loc[:, ["hod", "sum_by_hod", "prop_by_hod"]]
sum_hod
```

|    | hod | sum_by_hod | prop_by_hod |
|----|-----|------------|-------------|
| 0  | 0   | 20072      | 0.037992    |
| 1  | 1   | 20248      | 0.038325    |
| 2  | 2   | 18806      | 0.035595    |
| 3  | 3   | 19532      | 0.036969    |
| 4  | 4   | 20069      | 0.037986    |
| 5  | 5   | 21883      | 0.041419    |
| 6  | 6   | 23536      | 0.044548    |
| 7  | 7   | 21619      | 0.040920    |
| 8  | 8   | 21713      | 0.041098    |
| 9  | 9   | 22223      | 0.042063    |
| 10 | 10  | 24093      | 0.045602    |
| 11 | 11  | 23627      | 0.044720    |
| 12 | 12  | 23172      | 0.043859    |
| 13 | 13  | 23058      | 0.043643    |
| 14 | 14  | 22786      | 0.043129    |
| 15 | 15  | 23047      | 0.043623    |
| 16 | 16  | 23622      | 0.044711    |
| 17 | 17  | 23395      | 0.044281    |
| 18 | 18  | 24093      | 0.045602    |
| 19 | 19  | 22681      | 0.042930    |
| 20 | 20  | 22702      | 0.042970    |
| 21 | 21  | 20813      | 0.039394    |

|    |    |       |          |
|----|----|-------|----------|
| 22 | 22 | 20298 | 0.038419 |
| 23 | 23 | 21240 | 0.040202 |

Again, we can use `df.merge()` to concatenate new columns in `sum_hod` to the `data_grp2` table.

```
data_grp3 = pd.merge(data_grp2, sum_hod)
data_grp3
```

|       | hod | cod | freq_by_hodcod | sum_by_cod | prop_by_hodcod | sum_by_hod | \ |
|-------|-----|-----|----------------|------------|----------------|------------|---|
| 0     | 0   | A02 | 1              | 62         | 0.016129       | 20072      |   |
| 1     | 0   | A04 | 6              | 144        | 0.041667       | 20072      |   |
| 2     | 0   | A06 | 5              | 88         | 0.056818       | 20072      |   |
| 3     | 0   | A09 | 87             | 3111       | 0.027965       | 20072      |   |
| 4     | 0   | A15 | 7              | 209        | 0.033493       | 20072      |   |
| ...   | ... | ... | ...            | ...        | ...            | ...        |   |
| 16171 | 23  | N95 | 1              | 2          | 0.500000       | 21240      |   |
| 16172 | 23  | D24 | 1              | 1          | 1.000000       | 21240      |   |
| 16173 | 23  | N88 | 1              | 1          | 1.000000       | 21240      |   |
| 16174 | 23  | 067 | 1              | 1          | 1.000000       | 21240      |   |
| 16175 | 23  | V33 | 1              | 1          | 1.000000       | 21240      |   |

|       | prop_by_hod |
|-------|-------------|
| 0     | 0.037992    |
| 1     | 0.037992    |
| 2     | 0.037992    |
| 3     | 0.037992    |
| 4     | 0.037992    |
| ...   | ...         |
| 16171 | 0.040202    |
| 16172 | 0.040202    |
| 16173 | 0.040202    |
| 16174 | 0.040202    |
| 16175 | 0.040202    |

[16176 rows x 7 columns]

Check if we got the same result as the tidy paper.

```
data_grp3.query("cod in ['B16', 'E84', 'I21'] and hod == 8")
```

|      | hod | cod | freq_by_hodcod | sum_by_cod | prop_by_hodcod | sum_by_hod | \ |
|------|-----|-----|----------------|------------|----------------|------------|---|
| 5271 | 8   | B16 | 4              | 106        | 0.037736       | 21713      |   |
| 5397 | 8   | E84 | 3              | 111        | 0.027027       | 21713      |   |



```

5446      8  I21                2167      47510      0.045611      21713

      prop_by_hod
5271      0.041098
5397      0.041098
5446      0.041098

```

## 6.5 Deviation from the expected proportion

Finally, we can compute the deviation from the expected proportion (`prop_by_hod`)

```

data_grp3["diff_prop"] = (data_grp3["prop_by_hodcod"] -
data_grp3["prop_by_hod"])**2
data_grp3.head()

```

```

      hod  cod  freq_by_hodcod  sum_by_cod  prop_by_hodcod  sum_by_hod
prop_by_hod \
0      0  A02                1          62      0.016129      20072
0.037992
1      0  A04                6         144      0.041667      20072
0.037992
2      0  A06                5          88      0.056818      20072
0.037992
3      0  A09               87        3111      0.027965      20072
0.037992
4      0  A15                7         209      0.033493      20072
0.037992

      diff_prop
0      0.000478
1      0.000014
2      0.000354
3      0.000101
4      0.000020

```

And we can follow the same process as the tidy paper to get the distance from the expected proportion.

```

devi = data_grp3.groupby(["cod"]).agg(n=("freq_by_hodcod", lambda x:
x.sum()),
                                     dist=("diff_prop", lambda x:
x.mean()))
devi = devi.query("n > 50")
devi

```


```
      cod      n      dist
0     A01     51 0.000958
1     A02     62 0.000733
3     A04    144 0.000185
5     A06     88 0.000360
8     A09   3111 0.000030
...    ...    ...    ...
1172  Y33     60 0.000627
1173  Y34    780 0.000068
1183  Y57    111 0.000284
1190  Y83    174 0.000203
1193  Y86    363 0.000094
```

```
[450 rows x 3 columns]
```

Now let's do a visualization!


```
from plotnine import *
```

```
ggplot(devi, aes(x = 'n', y = 'dist')) + geom_point()
```



```
<ggplot: (401325638)>
```

```
ggplot(devi, aes(x = 'n', y = 'dist')) + geom_point()+\
scale_x_log10() +\
scale_y_log10() +\
geom_smooth(method = "lm")
```



```
<ggplot: (401161954)>
```

## 6.6 Use a residual plot to find unusual temporal patterns

We can use how well the data fit a linear model to find unusual temporal patterns. We need `statsmodels` to fit a linear model. As we learned from the previous visualization, the data are linearly related when the variables are log-transformed. Hence, we can directly fit a linear model on the log-transformed data.

```
# fit a linear model
import statsmodels.api as sm
import numpy as np
# log transformation
devi["log_n"] = np.log(devi["n"])
devi["log_dist"] = np.log(devi["dist"])
```

```
# fit a linear model
model = sm.OLS.from_formula("log_dist ~ log_n", data=devi)
result = model.fit()
result.summary()
```

### OLS Regression Results

|                   |                  |                     |           |       |        |        |
|-------------------|------------------|---------------------|-----------|-------|--------|--------|
| Dep. Variable:    | log_dist         | R-squared:          | 0.751     |       |        |        |
| Model:            | OLS              | Adj. R-squared:     | 0.750     |       |        |        |
| Method:           | Least Squares    | F-statistic:        | 1348.     |       |        |        |
| Date:             | Mon, 20 Feb 2023 | Prob (F-statistic): | 3.74e-137 |       |        |        |
| Time:             | 17:54:58         | Log-Likelihood:     | -436.72   |       |        |        |
| No. Observations: | 450              | AIC:                | 877.4     |       |        |        |
| Df Residuals:     | 448              | BIC:                | 885.7     |       |        |        |
| Df Model:         | 1                |                     |           |       |        |        |
| Covariance Type:  | nonrobust        |                     |           |       |        |        |
|                   | coef             | std err             | t         | P> t  | [0.025 | 0.975] |
| Intercept         | -4.0754          | 0.130               | -31.270   | 0.000 | -4.332 | -3.819 |
| log_n             | -0.8050          | 0.022               | -36.712   | 0.000 | -0.848 | -0.762 |
| Omnibus:          | 243.413          | Durbin-Watson:      | 1.160     |       |        |        |
| Prob(Omnibus):    | 0.000            | Jarque-Bera (JB):   | 1486.742  |       |        |        |
| Skew:             | 2.343            | Prob(JB):           | 0.00      |       |        |        |
| Kurtosis:         | 10.572           | Cond. No.           | 26.4      |       |        |        |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Put the fitted residuals into the `devi` dataframe.

```
# get the residual
devi["resid"] = result.resid
devi
```

```

    cod    n    dist    log_n    log_dist    resid
0   A01   51  0.000958  3.931826 -6.950290  0.290330
1   A02   62  0.000733  4.127134 -7.218939  0.178910
3   A04  144  0.000185  4.969813 -8.596625 -0.520395
```

```

5      A06      88  0.000360  4.477337  -7.929282 -0.249510
8      A09     3111 0.000030  8.042699 -10.422575  0.127419
...    ...     ...    ...         ...         ...         ...
1172   Y33      60  0.000627  4.094345  -7.373953 -0.002500
1173   Y34     780  0.000068  6.659294  -9.596346 -0.160034
1183   Y57     111  0.000284  4.709530  -8.165902 -0.299207
1190   Y83     174  0.000203  5.159055  -8.501431 -0.272856
1193   Y86     363  0.000094  5.894403  -9.270435 -0.449883

```


[450 rows x 6 columns]

Visualization: Use an arbitrary threshold, 1.5, to identify unusual temporal patterns.

```

ggplot(devi, aes(x = 'log_n', y = 'resid')) + geom_point() +
geom_hline(yintercept = 1.5)

```



```
<ggplot: (710113250)>
```

## 6.7 Visualize the unusual temporal patterns

Map the `cod` code to the full name of the cause of death.

```
map_cod = pd.read_csv("tidy_X_map_cod.csv")
map_cod.columns = ["cod", "cod_name"]
map_cod
```

```

      cod                                cod_name
0    A00                                Cholera
1    A01          Typhoid and paratyphoid fevers
2    A02          Other salmonella infections
3    A03                                Shigellosis
4    A04          Other bacterial intestinal infections
...    ...                                ...
1853 Y85          Sequelae of transport accidents
1854 Y86          Sequelae of other accidents
1855 Y87  Sequelae of intentional self harm, assault, an...
1856 Y88  Sequelae with surgical and medical care as ext...
1857 Y89          Sequelae of other external causes

[1858 rows x 2 columns]
```

```
devi2 = pd.merge(devi, map_cod)
devi2.head()
```

```

      cod    n    dist    log_n    log_dist    resid \
0  A01    51  0.000958  3.931826  -6.950290  0.290330
1  A02    62  0.000733  4.127134  -7.218939  0.178910
2  A04   144  0.000185  4.969813  -8.596625 -0.520395
3  A06    88  0.000360  4.477337  -7.929282 -0.249510
4  A09  3111  0.000030  8.042699 -10.422575  0.127419

                                cod_name
0          Typhoid and paratyphoid fevers
1          Other salmonella infections
2          Other bacterial intestinal infections
3                                Amebiasis
4  Diarrhea and gastroenteritis of infectious origin
```

Put the fitted information into the hour-level data, `data_grp3`.

```
data_grp3.head()
```

|             | hod | cod | freq_by_hodcod | sum_by_cod | prop_by_hodcod | sum_by_hod |
|-------------|-----|-----|----------------|------------|----------------|------------|
| prop_by_hod |     | \   |                |            |                |            |
| 0           | 0   | A02 | 1              | 62         | 0.016129       | 20072      |
|             |     |     |                |            |                | 0.037992   |
| 1           | 0   | A04 | 6              | 144        | 0.041667       | 20072      |
|             |     |     |                |            |                | 0.037992   |
| 2           | 0   | A06 | 5              | 88         | 0.056818       | 20072      |
|             |     |     |                |            |                | 0.037992   |
| 3           | 0   | A09 | 87             | 3111       | 0.027965       | 20072      |
|             |     |     |                |            |                | 0.037992   |
| 4           | 0   | A15 | 7              | 209        | 0.033493       | 20072      |
|             |     |     |                |            |                | 0.037992   |

|   | diff_prop |
|---|-----------|
| 0 | 0.000478  |
| 1 | 0.000014  |
| 2 | 0.000354  |
| 3 | 0.000101  |
| 4 | 0.000020  |

```
data_grp4 = pd.merge(data_grp3, devi2, on=["cod"])
data_grp4.head(10)
```

|             | hod | cod | freq_by_hodcod | sum_by_cod | prop_by_hodcod | sum_by_hod |
|-------------|-----|-----|----------------|------------|----------------|------------|
| prop_by_hod |     | \   |                |            |                |            |
| 0           | 0   | A02 | 1              | 62         | 0.016129       | 20072      |
|             |     |     |                |            |                | 0.037992   |
| 1           | 1   | A02 | 3              | 62         | 0.048387       | 20248      |
|             |     |     |                |            |                | 0.038325   |
| 2           | 2   | A02 | 9              | 62         | 0.145161       | 18806      |
|             |     |     |                |            |                | 0.035595   |
| 3           | 3   | A02 | 1              | 62         | 0.016129       | 19532      |
|             |     |     |                |            |                | 0.036969   |
| 4           | 4   | A02 | 3              | 62         | 0.048387       | 20069      |
|             |     |     |                |            |                | 0.037986   |
| 5           | 5   | A02 | 2              | 62         | 0.032258       | 21883      |
|             |     |     |                |            |                | 0.041419   |
| 6           | 6   | A02 | 2              | 62         | 0.032258       | 23536      |
|             |     |     |                |            |                | 0.044548   |
| 7           | 7   | A02 | 2              | 62         | 0.032258       | 21619      |
|             |     |     |                |            |                | 0.040920   |
| 8           | 8   | A02 | 1              | 62         | 0.016129       | 21713      |
|             |     |     |                |            |                | 0.041098   |
| 9           | 9   | A02 | 3              | 62         | 0.048387       | 22223      |
|             |     |     |                |            |                | 0.042063   |

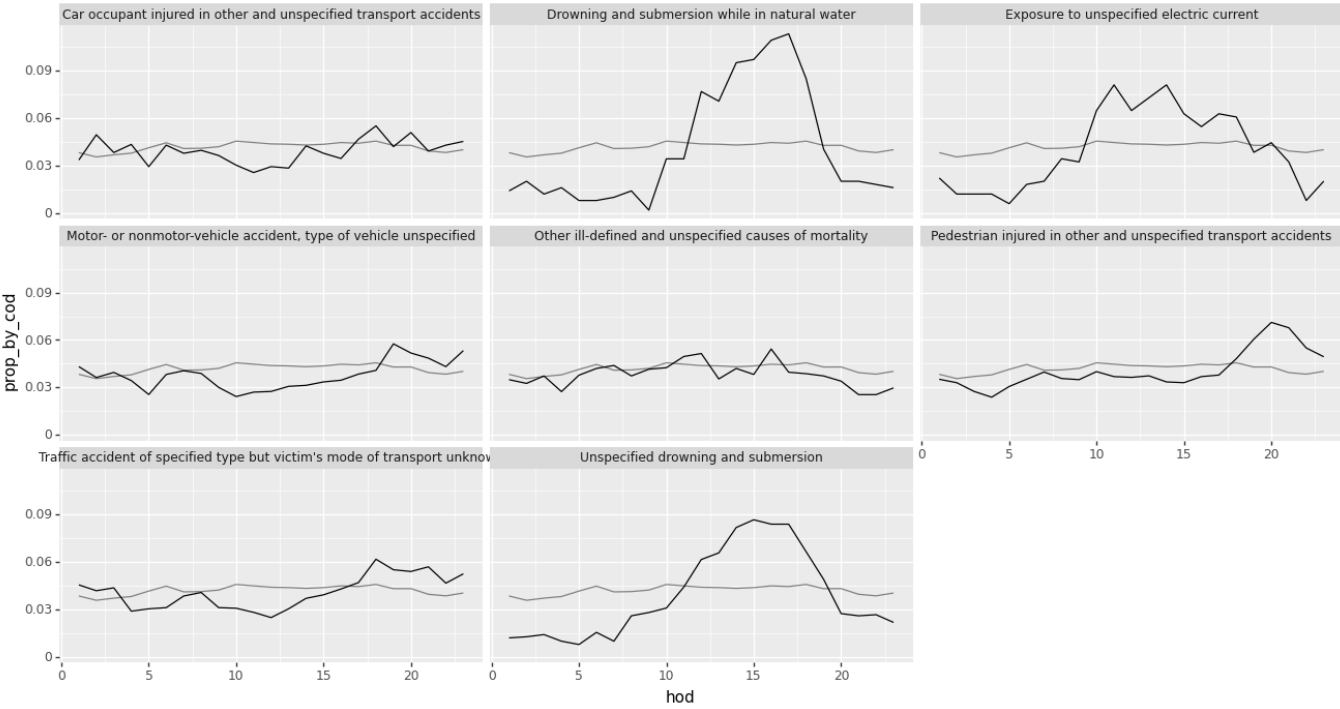
|   | diff_prop | n  | dist     | log_n    | log_dist  | resid   | \ |
|---|-----------|----|----------|----------|-----------|---------|---|
| 0 | 0.000478  | 62 | 0.000733 | 4.127134 | -7.218939 | 0.17891 |   |
| 1 | 0.000101  | 62 | 0.000733 | 4.127134 | -7.218939 | 0.17891 |   |
| 2 | 0.012005  | 62 | 0.000733 | 4.127134 | -7.218939 | 0.17891 |   |
| 3 | 0.000434  | 62 | 0.000733 | 4.127134 | -7.218939 | 0.17891 |   |
| 4 | 0.000108  | 62 | 0.000733 | 4.127134 | -7.218939 | 0.17891 |   |
| 5 | 0.000084  | 62 | 0.000733 | 4.127134 | -7.218939 | 0.17891 |   |
| 6 | 0.000151  | 62 | 0.000733 | 4.127134 | -7.218939 | 0.17891 |   |
| 7 | 0.000075  | 62 | 0.000733 | 4.127134 | -7.218939 | 0.17891 |   |
| 8 | 0.000623  | 62 | 0.000733 | 4.127134 | -7.218939 | 0.17891 |   |
| 9 | 0.000040  | 62 | 0.000733 | 4.127134 | -7.218939 | 0.17891 |   |

|   | cod_name                    |
|---|-----------------------------|
| 0 | Other salmonella infections |
| 1 | Other salmonella infections |
| 2 | Other salmonella infections |
| 3 | Other salmonella infections |
| 4 | Other salmonella infections |
| 5 | Other salmonella infections |
| 6 | Other salmonella infections |
| 7 | Other salmonella infections |
| 8 | Other salmonella infections |
| 9 | Other salmonella infections |

## Visualization

```
data_vis = data_grp4.query("n > 350 and (hod > 0 and hod < 24) and resid > 1.5").iloc[:184]
# set figure size
ggplot(data_vis, aes(x = 'hod', y = 'prop_by_hodcod')) +\
  geom_line(aes(y = "prop_by_hod"), colour = "grey") +\
  geom_line() +\
  facet_wrap('~cod_name', ncol = 3) +\
  theme(figure_size=(16, 8))
```





<ggplot: (707506116)>