

# Linear Regression

---

This course materials are based on the online material by Andrew Ng

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import sklearn
import plotly
```

## Linear Regression with one variable

Data overview - Housing price


```
house_size = [1000, 1250, 1400, 1500, 1560]
house_price_k = [200, 250, 370, 350, 400]
data_house = pd.DataFrame({'house_size': house_size,
                           'house_price_k': house_price_k})
data_house
```

```
.dataframe tbody tr th {
    vertical-align: top;
}

.dataframe thead th {
    text-align: right;
}
```

	house_size	house_price_k
0	1000	200
1	1250	250
2	1400	370
3	1500	350
4	1560	400

```
plt.scatter(data_house['house_size'], data_house['house_price_k'])
plt.xlabel('House Size')
plt.ylabel('House Price (K)')
plt.title('House Price vs House Size')
plt.grid(True)
# axis limit
plt.xlim(0, 2000)
plt.ylim(0, 500)
plt.show()
```



## Hypothesis


We will apply a supervised learning algorithm (i.e., linear regression in this case) to propose a hypothesis  $h$  with a given dataset. The hypothesis  $h$  determines an estimated "House Price (k)" ( $y$ ) based on a given "House Size" ( $x$ ). In short, The hypothesis  $h$  is a function of  $x$  that maps  $x$  to  $y$ .

$$y = h(x) = \beta_0 + \beta_1 x$$


We can try to plug in different values of  $\beta_0$  and  $\beta_1$  to see how the hypothesis  $h$  line looks like. To simplify the code, we need to define a plotting function that takes three parameters: dataset,  $\beta_0$ , and  $\beta_1$

```
def plot_house(data, b0, b1):
    # scatter plot
    plt.scatter(data['house_size'], data['house_price_k'])
    # hypothesis line
    x_series = np.linspace(0, 2000, 100)
    y_series = b0 + b1 * x_series
    plt.plot(x_series, y_series, 'r')
    # plot labels
    plt.title('House Price vs House Size')
    plt.xlabel('House Size')
    plt.ylabel('House Price (K)')
    # axis limit
    plt.xlim(0, 2000)
    plt.ylim(0, 500)
    # show plot
    plt.grid(True)
    plt.show()
```


```
plot_house(data=data_house, b0=0, b1=1)
```



```
plot_house(data=data_house, b0=200, b1=0)
```



```
plot_house(data=data_house, b0=-400, b1=0.5)
```



### Loss function

To evaluate the hypothesis  $h$ , we need to define a loss function. The loss function measures the difference between the estimated value ( $h(x)$ ) and the actual value ( $y$ ). The loss function is defined as

```
def loss_func(data, b0, b1):
    # calculate loss
    loss = 0
    for i in range(len(data)):
        x = data.loc[i, 'house_size']
        y = data.loc[i, 'house_price_k']
        loss += (y - (b0 + b1 * x)) ** 2
    return loss / len(data)
```

And you can check the loss function by plugging in different values of  $\beta_0$  and  $\beta_1$ .

```
loss_func(data_house, b0=0, b1=1)
```

```
1073800.0
```

```
loss_func(data_house, b0=200, b1=0)
```

```
18780.0
```

```
loss_func(data_house, b0=-400, b1=0.5)
```

3185.0

We can generate a loss surface by plotting the loss function with respect to  $\beta_0$  and  $\beta_1$ . The loss surface is a 3D plot that shows the loss function with respect to  $\beta_0$  and  $\beta_1$ . The loss surface is shown below.

```
# generate loss surface
b0_series = np.linspace(-1000, 1000, 100)
b1_series = np.linspace(-1, 1, 100)
b0_grid, b1_grid = np.meshgrid(b0_series, b1_series)
loss_grid = np.zeros((len(b0_series), len(b1_series)))
for i in range(len(b0_series)):
    for j in range(len(b1_series)):
        loss_grid[i, j] = loss_func(data=data_house,
                                     b0=b0_series[i],
                                     b1=b1_series[j])


# plot loss surface
import plotly.graph_objs as go
trace = go.Surface(x=b0_series, y=b1_series, z=loss_grid.T)
layout = go.Layout(
    title='Loss Surface',
    scene=dict(
        xaxis=dict(title='b0'),
        yaxis=dict(title='b1'),
        zaxis=dict(title='SSE'),
        width=900, height=900,
        margin=dict(r=20, l=10, b=10, t=50))
fig = go.Figure(data=[trace], layout=layout)
fig.show()
```

Let's try  $\beta$  s that seems to be the minimum of the loss surface.

```
loss_func(data_house, b0=-150, b1=0.35)
```

589.4500000000007

```
plot_house(data=data_house, b0=-150, b1=0.35)
```



### Matrix representation

So how do we calculate the best hypothesis  $h$  that can minimize the loss? The hypothesis  $h$  can be written in a matrix form as follows:

$$\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix}$$

$$\begin{bmatrix} 1 & x_1 & 1 & x_2 & \vdots & \vdots & 1 & x_m \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix} + \begin{bmatrix} \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_m \end{bmatrix}$$

where  $\epsilon$  is the error term. The matrix form of the hypothesis  $h$  is as follows:

$$\mathbf{y} = \mathbf{X} \mathbf{\beta} + \mathbf{\epsilon}$$

where

where  $\mathbf{y}$  is a  $m \times 1$  vector,  $\mathbf{X}$  is a  $m \times 2$  matrix,  $\mathbf{\beta}$  is a  $2 \times 1$  vector, and  $\mathbf{\epsilon}$  is a  $m \times 1$  vector.

In our dataset, the formula should look like this:

$$\begin{bmatrix} 200 \\ 250 \\ 370 \\ 350 \\ 400 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 1000 & 1 & 1250 & 1 & 1400 & 1 & 1500 & 1 & 1560 \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix} + \begin{bmatrix} \epsilon_1 \\ \epsilon_2 \\ \epsilon_3 \\ \epsilon_4 \\ \epsilon_5 \end{bmatrix}$$

This is equivalent to the following linear equations:

$$\begin{aligned} 200 &= \beta_0 + 1000 \beta_1 + \epsilon_1 \\ 250 &= \beta_0 + 1250 \beta_1 + \epsilon_2 \\ 370 &= \beta_0 + 1400 \beta_1 + \epsilon_3 \\ 350 &= \beta_0 + 1500 \beta_1 + \epsilon_4 \\ 400 &= \beta_0 + 1560 \beta_1 + \epsilon_5 \end{aligned}$$

### Ordinary Least Squares

The Ordinary Least Squares (OLS) is a method to estimate the parameters  $\beta_0$  and  $\beta_1$  in the hypothesis  $h$  by minimizing the sum of squared errors (SSE). The SSE is defined as follows:

$$\text{SSE} = \sum_{i=1}^m (y_i - h(x_i))^2$$

Or in matrix form:

$$\text{SSE} = (\mathbf{y} - \mathbf{X}\mathbf{\beta})^T (\mathbf{y} - \mathbf{X}\mathbf{\beta})$$

The minimization problem can be extended as

$$\begin{aligned} \arg \min_{\mathbf{\beta}} \text{SSE} &= \arg \min_{\mathbf{\beta}} f(\mathbf{\beta}) \quad \&= \arg \min_{\mathbf{\beta}} (\mathbf{y} - \mathbf{X}\mathbf{\beta})^T (\mathbf{y} - \mathbf{X}\mathbf{\beta}) \\ &\quad \&= \arg \min_{\mathbf{\beta}} (\mathbf{y}^T - \mathbf{\beta}^T \mathbf{X}^T) (\mathbf{y} - \mathbf{X}\mathbf{\beta}) \\ &\quad \&= \arg \min_{\mathbf{\beta}} \mathbf{y}^T \mathbf{y} - \mathbf{y}^T \mathbf{X} \mathbf{\beta} - \mathbf{\beta}^T \mathbf{X}^T \mathbf{y} + \mathbf{\beta}^T \mathbf{X}^T \mathbf{X} \mathbf{\beta} \\ &\quad \&= \arg \min_{\mathbf{\beta}} \mathbf{y}^T \mathbf{y} - 2 \mathbf{y}^T \mathbf{X} \mathbf{\beta} + \mathbf{\beta}^T \mathbf{X}^T \mathbf{X} \mathbf{\beta} \end{aligned}$$

Find the partial derivatives of  $f(\mathbf{\beta})$  with respect to  $\mathbf{\beta}$  and set it to zero:

$$\frac{\partial f}{\partial \mathbf{\beta}} = -\mathbf{X}^T \mathbf{y} + \mathbf{X}^T \mathbf{X} \mathbf{\beta} = 0$$

$$\mathbf{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Solve the OLS problem on our dataset

Now, we can plug in our dataset to this formula to obtain the optimum  $\mathbf{\beta}$ . First, we need to define  $\mathbf{X}$  and  $\mathbf{y}$  in Python.

```
X = data_house['house_size']
y = data_house['house_price_k']
X = np.array(X)
X = np.vstack([np.ones(len(X)), X]).T # add intercept
y = np.array(y)
# show
display("X:", X)
display("y:", y)
```

'X:'

```
array([[1.00e+00, 1.00e+03],
       [1.00e+00, 1.25e+03],
       [1.00e+00, 1.40e+03],
       [1.00e+00, 1.50e+03],
       [1.00e+00, 1.56e+03]])
```

'y:'

```
array([200, 250, 370, 350, 400])
```

Obtain the  $\mathbf{\beta}$  based on the formula:

$$\mathbf{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \tag{1}$$

```
XtX = np.dot(X.T, X)
XtXi = np.linalg.inv(XtX)
Xy = np.dot(X.T, y)
b = np.dot(XtXi, Xy)
print(b)
```

```
[-169.78139905    0.36049285]
```

We can wrap the computation as a solver function, which takes  $X$  and  $y$  as input parameters.

```
def solve_OLS(X, y):  
    XtX = np.dot(X.T, X)  
    XtXi = np.linalg.inv(XtX)  
    Xy = np.dot(X.T, y)  
    b = np.dot(XtXi, Xy)  
    return b
```

Should return the same  $\beta$  as the previous computation.

```
solve_OLS(X, y)
```


```
array([-169.78139905,    0.36049285])
```

Great! Now we have the optimal  $\beta_0 = -169.78$  and  $\beta_1 = 0.36$ . Let's validate this result.

```
loss_func(data_house, b0=b[0], b1=b[1])
```

```
552.5278219395883
```

```
plot_house(data=data_house, b0=b[0], b1=b[1])
```





## Sklearn - Linear Regression

The Python package `sklearn` provides a function to solve the OLS problem. It's noted that we don't need to add intercepts to the  $X$  matrix.

```
from sklearn.linear_model import LinearRegression
X = np.array(data_house['house_size']).reshape((-1, 1))
y = np.array(data_house['house_price_k'])
model = LinearRegression()
model.fit(X, y)
```

▼ LinearRegression

LinearRegression()

```
print("coef.: ", model.coef_)
print("intercept: ", model.intercept_)
```

```
coef.: [0.36049285]
intercept: -169.78139904610492
```

We can also use this model to predict a new data point.

```
new_x = [[750]] # predict house price for 750 sqft house
model.predict(new_x)
```

```
array([100.58823529])
```

## Multi-variable Linear Regression

### Medical insurance dataset

In this section, we will use a real-world dataset to demonstrate the multi-variable linear regression. The dataset contains information about medical insurance costs for 1338 people. The dataset is available on [Kaggle](#).

```
data = pd.read_csv("insurance.csv")
data
```

```
.dataframe tbody tr th {
    vertical-align: top;
}

.dataframe thead th {
    text-align: right;
}
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400

	age	sex	bmi	children	smoker	region	charges
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

### The regression problem

Say, we want to know how the listed factors, which include **age**, **sex**, **bmi**, **children**, **smoker**, **region**, affect the **charges** of medical insurance. We can define the hypothesis  $h$  as follows:

$$h(x) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \beta_4 x_4 + \beta_5 x_5 + \beta_6 x_6 + \beta_7 x_7 + \beta_8 x_8$$

where

- $x_1 = \text{age}$
- $x_2 = \text{sex}$ ; 0 for male, 1 for female
- $x_3 = \text{bmi}$
- $x_4 = \text{children}$
- $x_5 = \text{smoker}$ ; 0 for non-smoker (no), 1 for smoker (yes)
- $x_6 = \text{region}$  is **southwest** (1) or not (0)
- $x_7 = \text{region}$  is **southeast** (1) or not (0)
- $x_8 = \text{region}$  is **northwest** (1) or not (0)

Note that we turned a categorical variable **region** into three binary variables. This is called one-hot encoding.

### Solve the OLS problem

Now, let's try to solve this problem using the formula  $(1)$ . We need to define the matrix  $X$  and  $y$  before we can compute the  $\beta$ .

```
X = data.iloc[:, :-1]
y = data["charges"]
display(X)
display(y)
```

```
.dataframe tbody tr th {
    vertical-align: top;
}

.dataframe thead th {
    text-align: right;
}
```

	age	sex	bmi	children	smoker	region
0	19	female	27.900	0	yes	southwest
1	18	male	33.770	1	no	southeast
2	28	male	33.000	3	no	southeast
3	33	male	22.705	0	no	northwest
4	32	male	28.880	0	no	northwest
...	...	...	...	...	...	...
1333	50	male	30.970	3	no	northwest
1334	18	female	31.920	0	no	northeast
1335	18	female	36.850	0	no	southeast
1336	21	female	25.800	0	no	southwest
1337	61	female	29.070	0	yes	northwest

1338 rows × 6 columns

```

0      16884.92400
1      1725.55230
2      4449.46200
3      21984.47061
4      3866.85520
...
1333   10600.54830
1334    2205.98080
1335    1629.83350
1336    2007.94500
1337    29141.36030
Name: charges, Length: 1338, dtype: float64

```

```

# turn categorical variable into dummy variables
X = pd.get_dummies(X, drop_first=True) # very handy function!
display(X)

```

```

.dataframe tbody tr th {
    vertical-align: top;
}

.dataframe thead th {
    text-align: right;
}

```

	age	bmi	children	sex_male	smoker_yes	region_northwest	region_southeast	region_southwest
0	19	27.900	0	0	1	0	0	1
1	18	33.770	1	1	0	0	1	0
2	28	33.000	3	1	0	0	1	0
3	33	22.705	0	1	0	1	0	0

	age	bmi	children	sex_male	smoker_yes	region_northwest	region_southeast	region_southwest
4	32	28.880	0	1	0	1	0	0
...	...	...	...	...	...	...	...	...
1333	50	30.970	3	1	0	1	0	0
1334	18	31.920	0	0	0	0	0	0
1335	18	36.850	0	0	0	0	1	0
1336	21	25.800	0	0	0	0	0	1
1337	61	29.070	0	0	1	1	0	0

1338 rows × 8 columns

Add intercept to the **X** matrix and turn both **X** and **y** into numpy arraies.

```
# add intercept
X = np.array(X)
X = np.hstack([np.ones((len(X), 1)), X])
y = np.array(y)
```

Check the dimensions of **X** and **y**.

```
print("X shape: ", X.shape)
display(X)
print("y shape: ", y.shape)
display(y)
```

X shape: (1338, 9)

```
array([[ 1. , 19. , 27.9 , ...,  0. ,  0. ,  1. ],
       [ 1. , 18. , 33.77, ...,  0. ,  1. ,  0. ],
       [ 1. , 28. , 33. , ...,  0. ,  1. ,  0. ],
       ...,
       [ 1. , 18. , 36.85, ...,  0. ,  1. ,  0. ],
       [ 1. , 21. , 25.8 , ...,  0. ,  0. ,  1. ],
       [ 1. , 61. , 29.07, ...,  1. ,  0. ,  0. ]])
```

y shape: (1338,)

```
array([16884.924 , 1725.5523, 4449.462 , ..., 1629.8335, 2007.945 ,
       29141.3603])
```

Plug in the **X** and **y** into the formula  $\beta(1)$  to obtain the  $\beta$ .

```
solve_OLS(X, y)
```

```
array([-11938.53857617,    256.85635254,    339.19345361,    475.50054515,
       -131.3143594 ,   23848.53454191,   -352.96389942,  -1035.02204939,
       -960.0509913 ])
```

We can solve the same problem using the `sklearn` library

```
X = data.iloc[:, :-1]
X = pd.get_dummies(X, drop_first=True)
y = data["charges"]
model = LinearRegression()
model.fit(X, y)
```

▼ LinearRegression

LinearRegression()

```
print("coef.: ", model.coef_)
print("intercept: ", model.intercept_)
```

```
coef.: [ 256.85635254  339.19345361  475.50054515 -131.3143594
 23848.53454191 -352.96389942 -1035.02204939 -960.0509913 ]
intercept: -11938.5385761715
```