

Lecture 9-1: Linear Regression

Dr. James Chen, Animal Data Scientist, School of Animal Sciences

2023 APSC-5984 SS: Agriculture Data Science




Let's Start With an Example - Housing Prices

X's: “input” variable / features / independent variables

Y's: “output” variable / labels / response variables

	X		Y
	house_size	house_price_k	
0	1000	200	
1	1250	250	
2	1400	370	
3	1500	350	
4	1560	400	
5	1900	???	

How do we predict the price?



	X		Y
	house_size	house_price_k	
0	1000		200
1	1250		250
2	1400		370
3	1500		350
4	1560		400

Training Set

Learning Algorithm

Hypothesis

X
(House size)

Y
(House price)

$$y = h(x) = \beta_0 + \beta_1 x$$

The hypothesis that maps x to y

e.g., **Linear regression** is a type of supervised learning algorithm:

(x, y) - one training example

feature label


$$y = h(x) = \beta_0 + \beta_1 x$$

The hypothesis that maps x to y (predicted)


Parameter β_0 : intercept

Parameter β_1 : slope

$$\beta: (\beta_0, \beta_1) = (0, 1)$$




$$\beta: (\beta_0, \beta_1) = (200, 0)$$



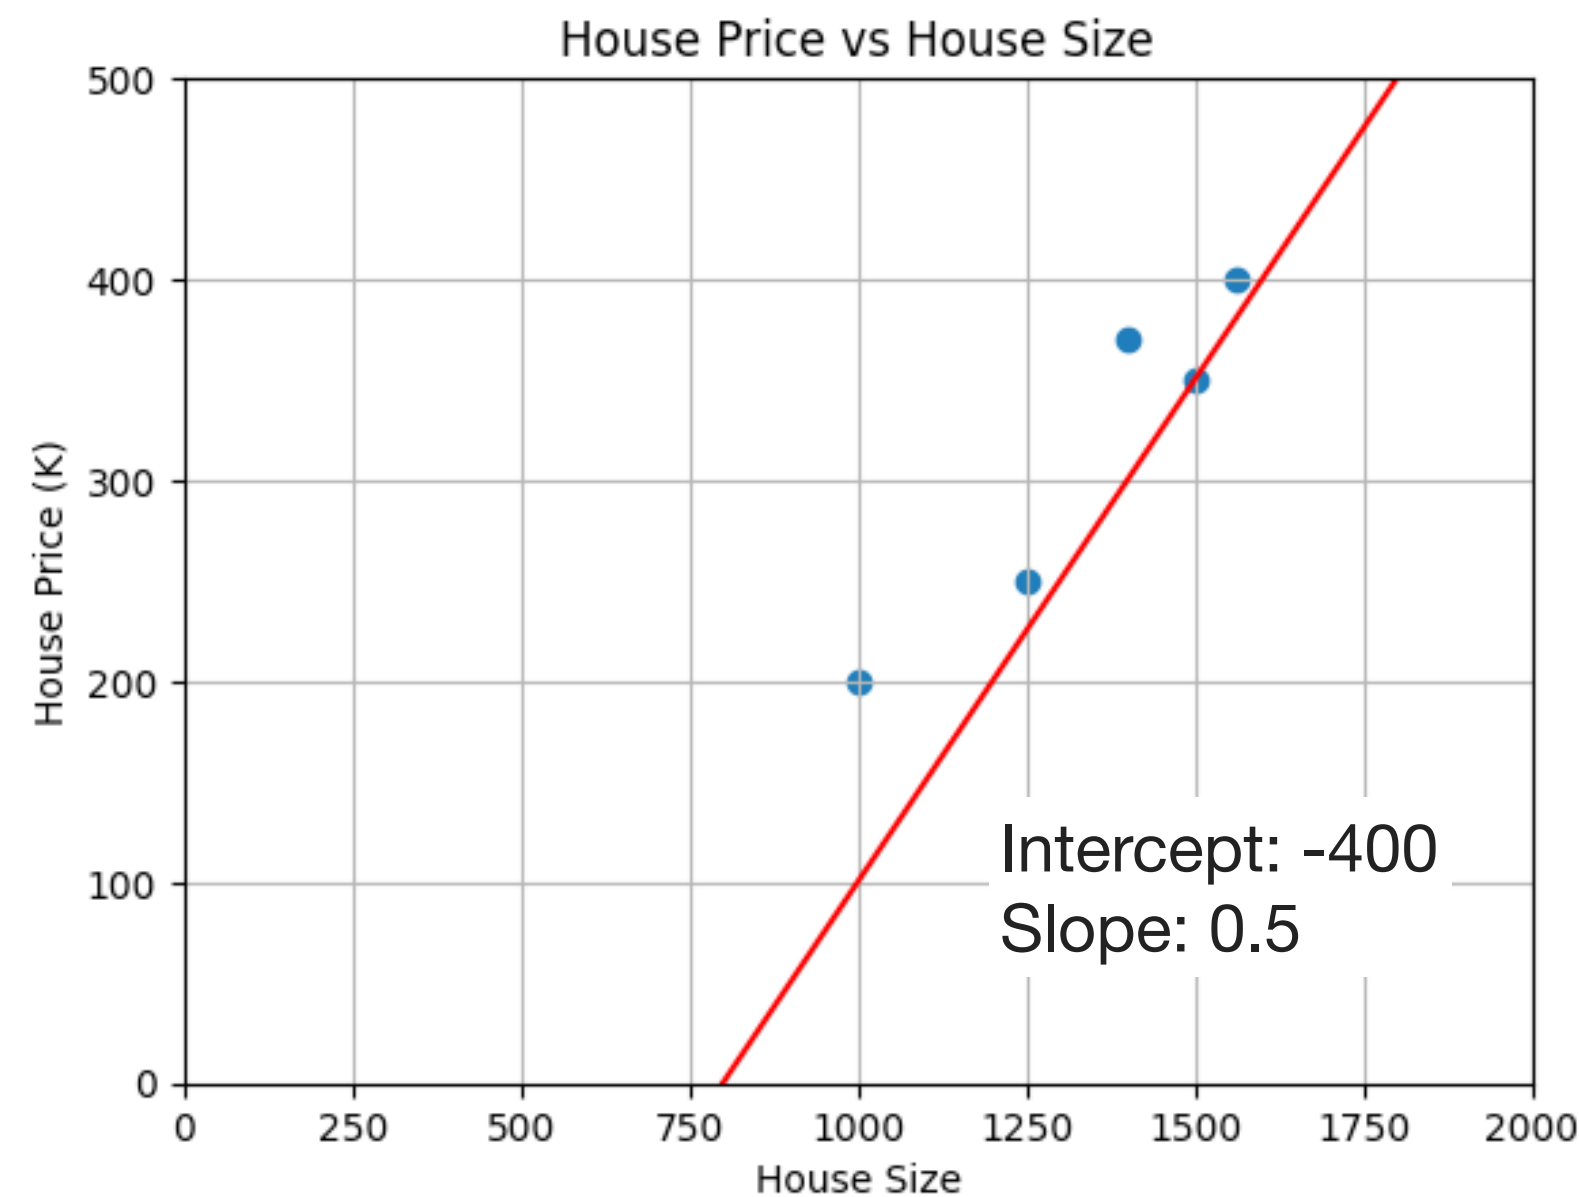
Is it the best hypothesis?

$$\beta: (\beta_0, \beta_1) = (-400, 0.5)$$



Is it the best hypothesis?

$$\beta: (\beta_0, \beta_1) = (-400, 0.5)$$



Let's define the loss function f : mean squared error (MSE)

MSE

$$f(\beta_0, \beta_1) = \frac{1}{n} \sum_{i=1}^n (\underbrace{h(x_i)}_{\text{Predicted } y} - \underbrace{y_i}_{\text{Observed } y})^2$$

Predicted y

$$h(x_i) = \underbrace{\beta_0}_{\text{Intercept}} + \underbrace{\beta_1 x_i}_{\text{Slope}}$$

Intercept

Slope

Loss Function

MSE


$$f(\beta_0, \beta_1) = \frac{1}{n} \sum_{i=1}^n (\underbrace{h(x_i)}_{\text{Predicted } y} - \underbrace{y_i}_{\text{Observed } y})^2$$

```
def loss_func(data, b0, b1):  
    # calculate loss  
    loss = 0  
    for i in range(len(data)):  
        x = data.loc[i, 'house_size']  
        y = data.loc[i, 'house_price_k']  
        loss += (y - (b0 + b1 * x)) ** 2  
    return loss / len(data)
```

```
loss_func(data_house, b0=0, b1=1)
```

1073800.0


Intercept: 0
Slope: 1



```
loss_func(data_house, b0=200, b1=0)
```

18780.0


Intercept: 200
Slope: 0



```
loss_func(data_house, b0=-400, b1=0.5)
```

3185.0

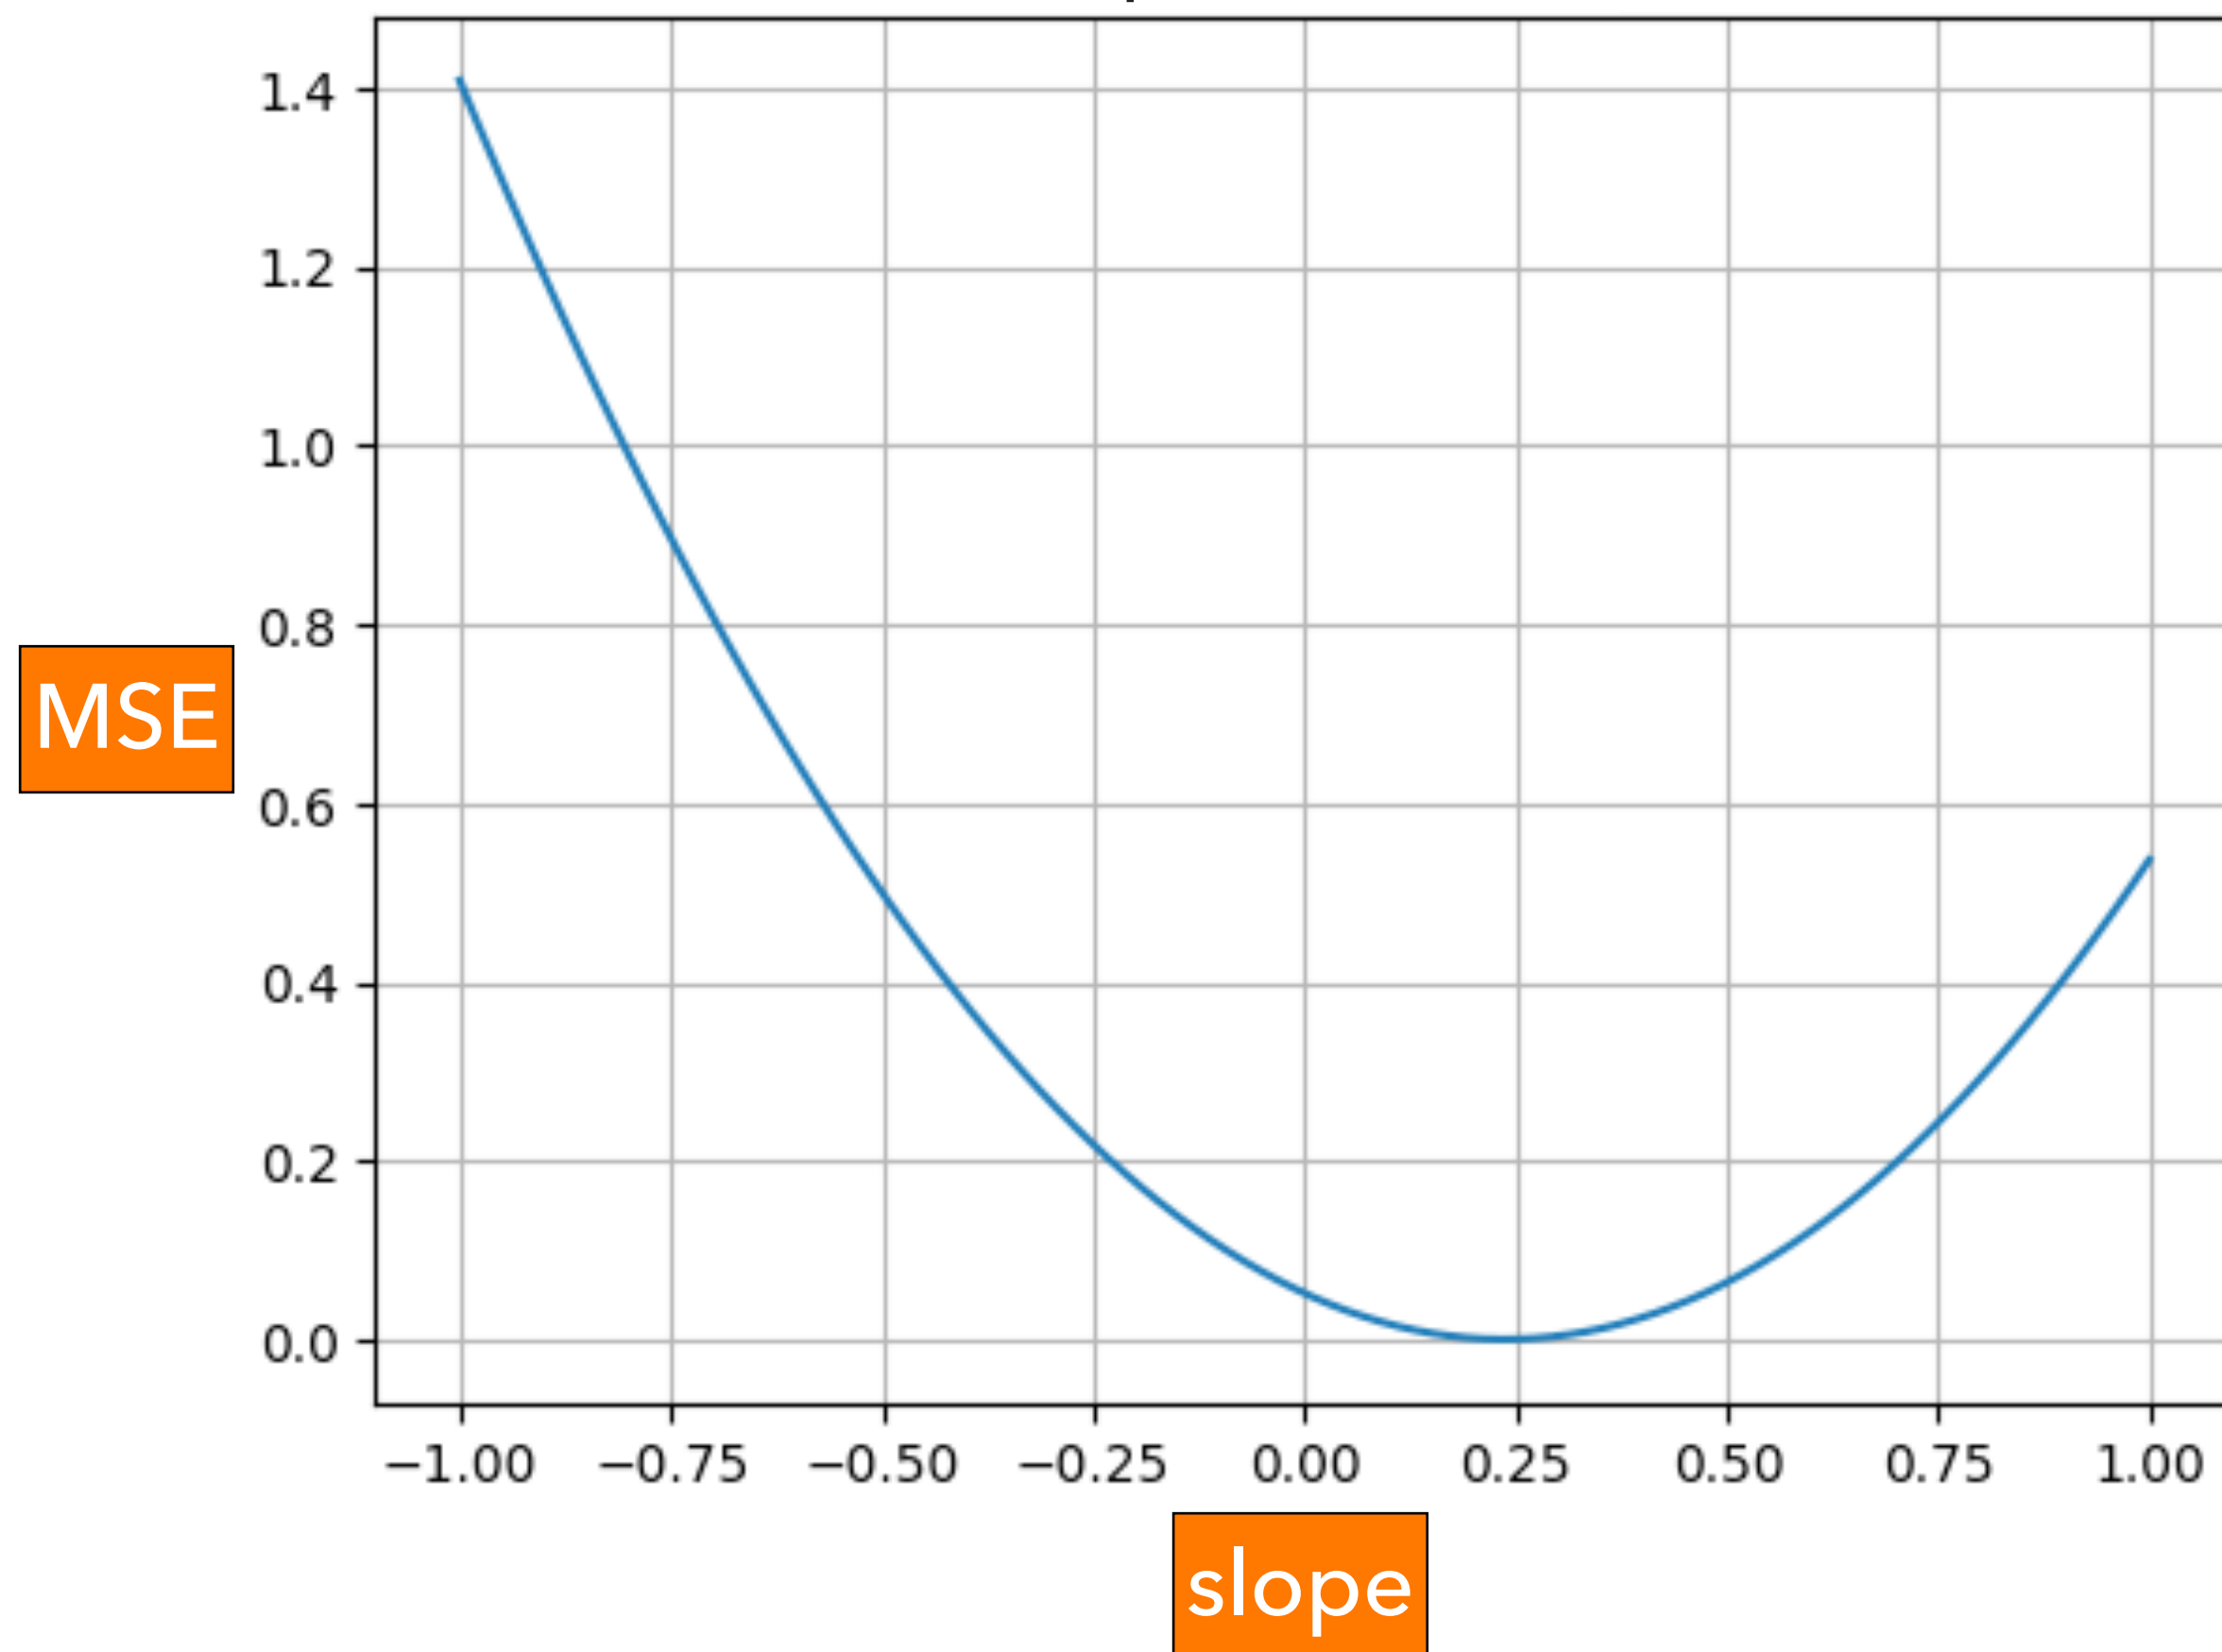
Intercept: -400
Slope: 0.5



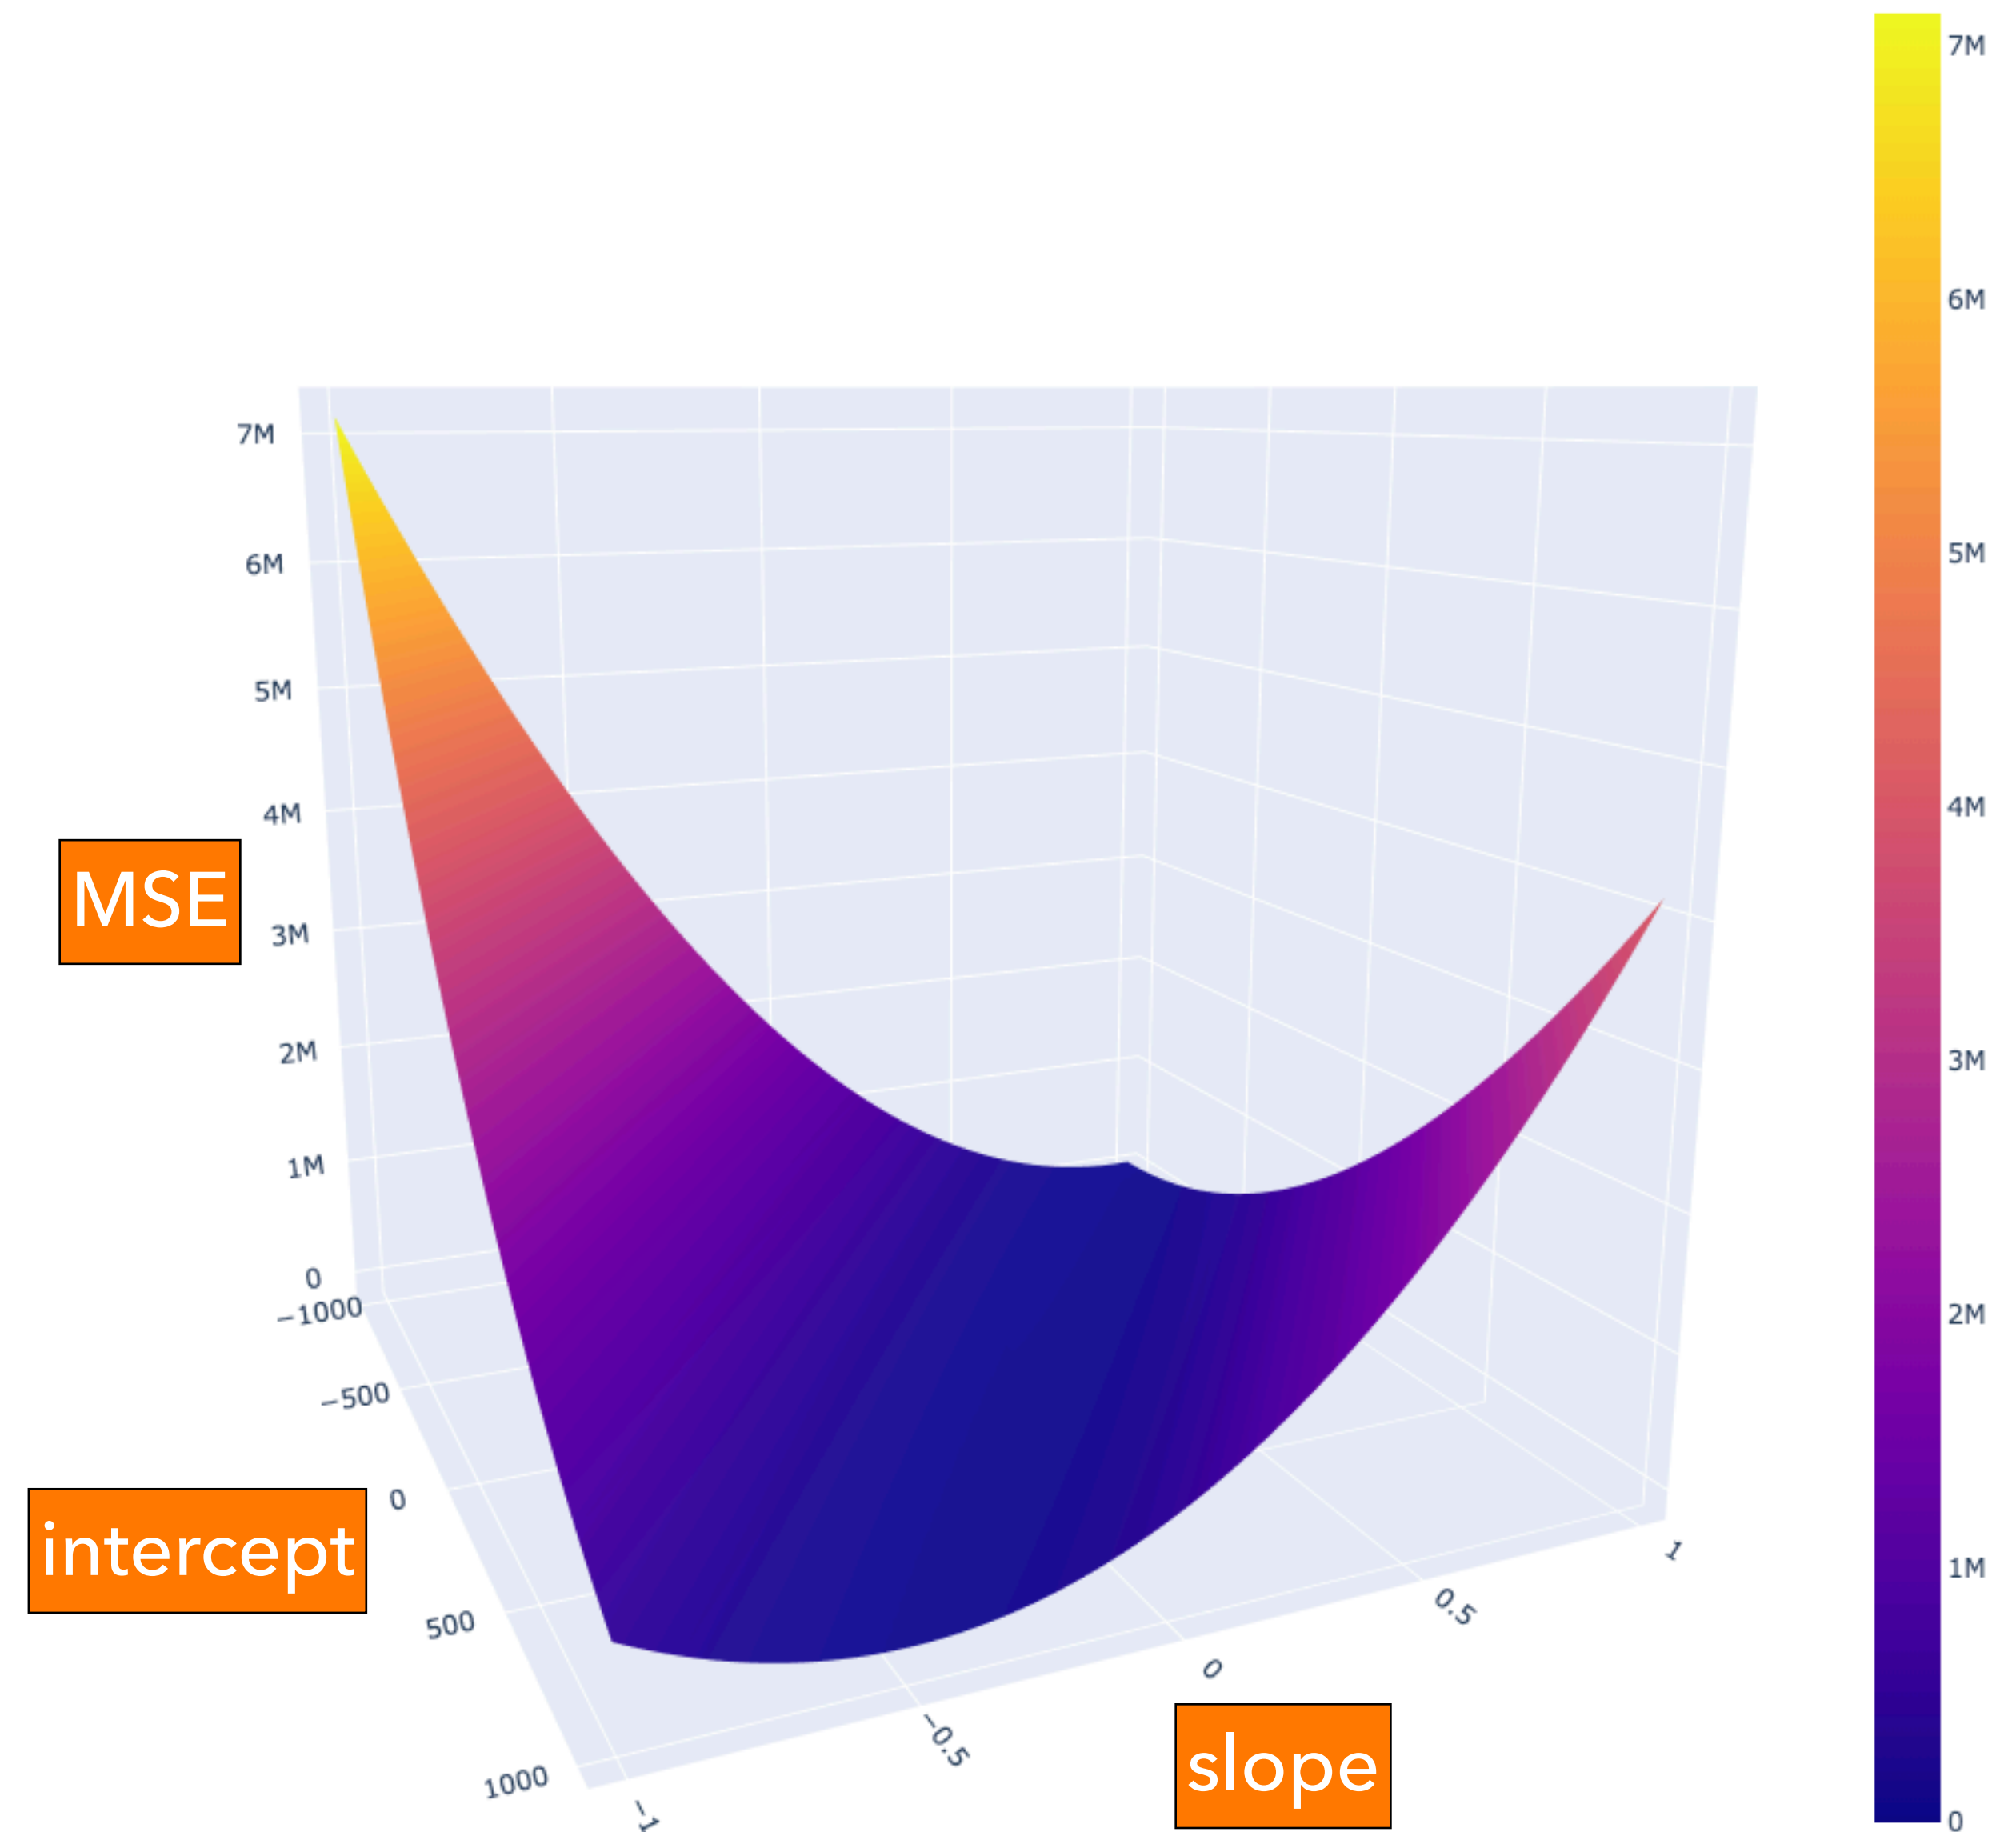
Loss Curve and Surface

How do we find a hypothesis that minimizes the loss? Differentiation!

We fixed the intercept to 0



Loss Surface



Matrix multiplication

Rule 1

$$\begin{bmatrix} a & b \\ c & d \\ e & f \end{bmatrix} \begin{bmatrix} g & h \\ i & j \end{bmatrix} = \begin{bmatrix} ag + bi & ah + bj \\ cg + di & ch + dj \\ eg + fi & eh + fj \end{bmatrix}$$

Rule 2

$$AB \neq BA$$

Matrix transpose

Rule 1

$$\begin{bmatrix} a & b \\ c & d \\ e & f \end{bmatrix}^T = \begin{bmatrix} a & c & e \\ b & d & f \end{bmatrix}$$

Rule 2

$$(A + B)^T = A^T + B^T$$

Rule 3

$$(AB)^T = B^T A^T$$

Rule 4

$$(ABC)^T = C^T B^T A^T$$

Matrix inverse

Rule 1

$$A = \begin{bmatrix} a & b \\ c & d \end{bmatrix}$$

$$A^{-1} = \frac{1}{ad - bc} \begin{bmatrix} d & -b \\ -c & a \end{bmatrix}$$

Rule 2

$$A^{-1}A = I$$

$$AA^{-1} = I$$

Matrix derivative

$$\begin{aligned} \frac{\partial \beta^T A \beta}{\partial \beta} &= 2\beta^T A \\ &= 2A\beta \end{aligned}$$

Data frame

Linear equations

Matrix representation

	X		Y
	house_size	house_price_k	
0	1000		200
1	1250		250
2	1400		370
3	1500		350
4	1560		400

$$\begin{aligned}200 &= \beta_0 + 1000\beta_1 + \epsilon_1 \\250 &= \beta_0 + 1250\beta_1 + \epsilon_2 \\370 &= \beta_0 + 1400\beta_1 + \epsilon_3 \\350 &= \beta_0 + 1500\beta_1 + \epsilon_4 \\400 &= \beta_0 + 1560\beta_1 + \epsilon_5\end{aligned}$$

$$\begin{bmatrix} 200 \\ 250 \\ 370 \\ 350 \\ 400 \end{bmatrix} = \begin{bmatrix} 1 & 1000 \\ 1 & 1250 \\ 1 & 1400 \\ 1 & 1500 \\ 1 & 1560 \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix} + \begin{bmatrix} \epsilon_1 \\ \epsilon_2 \\ \epsilon_3 \\ \epsilon_4 \\ \epsilon_5 \end{bmatrix}$$

We want to minimize the error term

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}$$

The Ordinary Least Squares (OLS) is a method to estimate the parameters β_0 and β_1 in the hypothesis h by minimizing the sum of squared errors (SSE). The SSE is defined as follows:

$$\text{SSE} = \sum_{i=1}^m (y_i - h(x_i))^2$$

Or in matrix form:

$$\text{SSE} = (\mathbf{y} - \mathbf{X}\beta)^T (\mathbf{y} - \mathbf{X}\beta)$$

MSE (Recap)

$$f(\beta_0, \beta_1) = \frac{1}{n} \sum_{i=1}^n (h(x_i) - y_i)^2$$

The minimization problem can be extended as

$$\begin{aligned} \arg \min_{\beta} \text{SSE} &= \arg \min_{\beta} f(\beta) = \arg \min_{\beta} (\mathbf{y} - \mathbf{X}\beta)^T (\mathbf{y} - \mathbf{X}\beta) \\ &= \arg \min_{\beta} (\mathbf{y}^T - \beta^T \mathbf{X}^T) (\mathbf{y} - \mathbf{X}\beta) \\ &= \arg \min_{\beta} \mathbf{y}^T \mathbf{y} - \mathbf{y}^T \mathbf{X}\beta - \beta^T \mathbf{X}^T \mathbf{y} + \beta^T \mathbf{X}^T \mathbf{X}\beta \\ &= \arg \min_{\beta} \mathbf{y}^T \mathbf{y} - 2\mathbf{y}^T \mathbf{X}\beta + \beta^T \mathbf{X}^T \mathbf{X}\beta \end{aligned}$$

Find the partial derivatives of $f(\beta)$ with respect to β and set it to zero:

$$\frac{\partial f}{\partial \beta} = \beta^T \mathbf{X}^T \mathbf{X} - \mathbf{X}^T \mathbf{y} = 0$$

$$\beta = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \quad \leftarrow \text{This is the learning algorithm!}$$

Now, we can plug in our dataset to this formula to obtain the optimum β . First, we need to define X and y in Python.

```
X = data_house['house_size']
y = data_house['house_price_k']
X = np.array(X)
X = np.vstack([np.ones(len(X)), X]).T # add intercept
y = np.array(y)
# show
display("X:", X)
display("y:", y)
```

Python representation

```
'X:'
array([[1.00e+00, 1.00e+03],
       [1.00e+00, 1.25e+03],
       [1.00e+00, 1.40e+03],
       [1.00e+00, 1.50e+03],
       [1.00e+00, 1.56e+03]])

'y:'
array([200, 250, 370, 350, 400])
```

Math representation

$$\begin{bmatrix} 200 \\ 250 \\ 370 \\ 350 \\ 400 \end{bmatrix} = \begin{bmatrix} 1 & 1000 \\ 1 & 1250 \\ 1 & 1400 \\ 1 & 1500 \\ 1 & 1560 \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix} + \begin{bmatrix} \epsilon_1 \\ \epsilon_2 \\ \epsilon_3 \\ \epsilon_4 \\ \epsilon_5 \end{bmatrix}$$

Manually Solve the OLS Problem

Obtain the β based on the formula:

$$\beta = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \quad (1)$$

<code>XtX = np.dot(X.T, X)</code>	→	Matrix multiplication	$\mathbf{X}^T \mathbf{X}$
<code>XtXi = np.linalg.inv(XtX)</code>	→	Matrix inverse	$(\mathbf{X}^T \mathbf{X})^{-1}$
<code>Xy = np.dot(X.T, y)</code>	→	Matrix multiplication	$\mathbf{X}^T \mathbf{y}$
<code>b = np.dot(XtXi, Xy)</code>	→	Matrix multiplication	$(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$
<code>print(b)</code>			

MagicPython

`[-169.78139905 0.36049285]`

Parameter β_0
intercept

Parameter β_1
slope

Wrap Up the Solver Into a Function

We can wrap the computation as a solver function, which takes X and y as input parameters.

```
def solve_OLS(X, y):  
    XtX = np.dot(X.T, X)  
    XtXi = np.linalg.inv(XtX)  
    Xy = np.dot(X.T, y)  
    b = np.dot(XtXi, Xy)  
    return b
```

MagicPython

Should return the same β as the previous computation.

```
solve_OLS(X, y)
```

MagicPython

```
array([-169.78139905,  0.36049285])
```

Parameter β_0
intercept

Parameter β_1
slope

Great! Now we have the optimal $\beta_0 = -169.78$ and $\beta_1 = 0.36$. Let's validate this result.


```
loss_func(data_house, b0=b[0], b1=b[1])
```

MagicPython

552.5278219395883

```
plot_house(data=data_house, b0=b[0], b1=b[1])
```

MagicPython



The Python package `sklearn` provides a function to solve the OLS problem. It's noted that we don't need to add intercepts to the X matrix.

```
from sklearn.linear_model import LinearRegression
X = np.array(data_house['house_size']).reshape((-1, 1))
y = np.array(data_house['house_price_k'])
model = LinearRegression()
model.fit(X, y)
```

MagicPython

▼ LinearRegression
LinearRegression()

```
print("coef.: ", model.coef_)
print("intercept: ", model.intercept_)
```

```
coef.: [0.36049285]
intercept: -169.78139904610492
```

We can also use this model to predict a new data point.

```
new_x = [[750]] # predict house price for 750 sqft house
model.predict(new_x)
```

```
array([100.58823529])
```